

Operating Systems

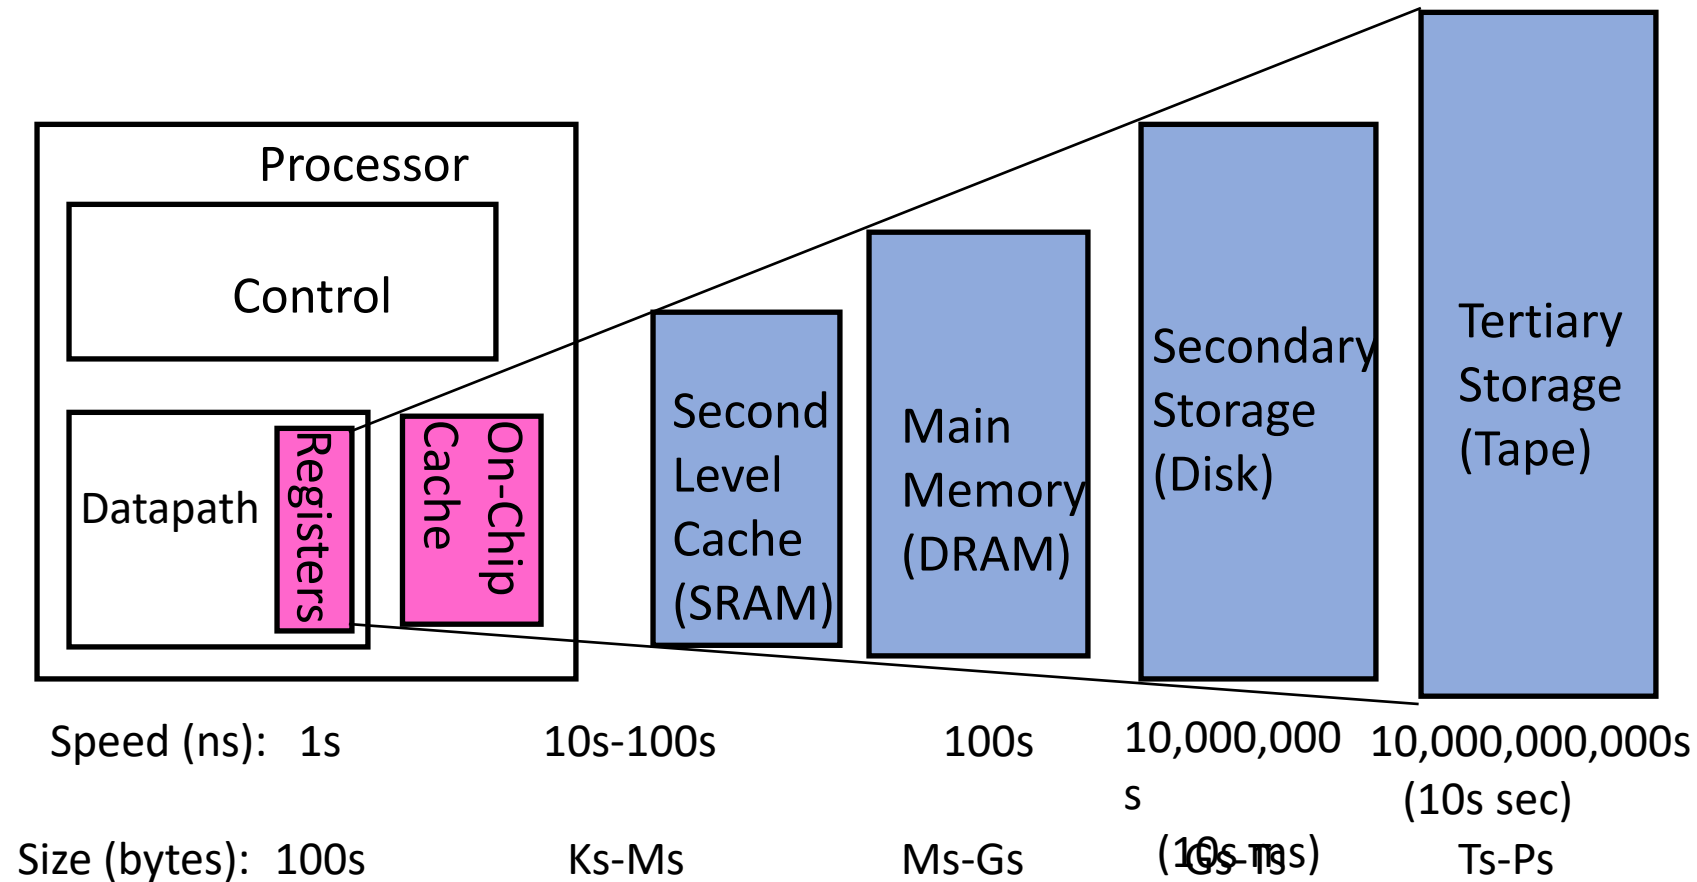
Lecture 8

Demand Paging

Prof. Mengwei Xu

Recap: Memory Hierarchy

- Speed, Size, and Cost: take advantage of each level




Recap: Locality

- Temporal locality (时间局部性): If at one point a particular memory location is referenced, then it is likely that the same location will be referenced again in the near future.
 - To leverage: keep recently accessed data items closer to processor
- Spatial locality (空间局部性): if a particular storage location is referenced at a particular time, then it is likely that nearby memory locations will be referenced in the near future.
 - Move contiguous blocks to the upper levels

Recap: TLB as a Cache

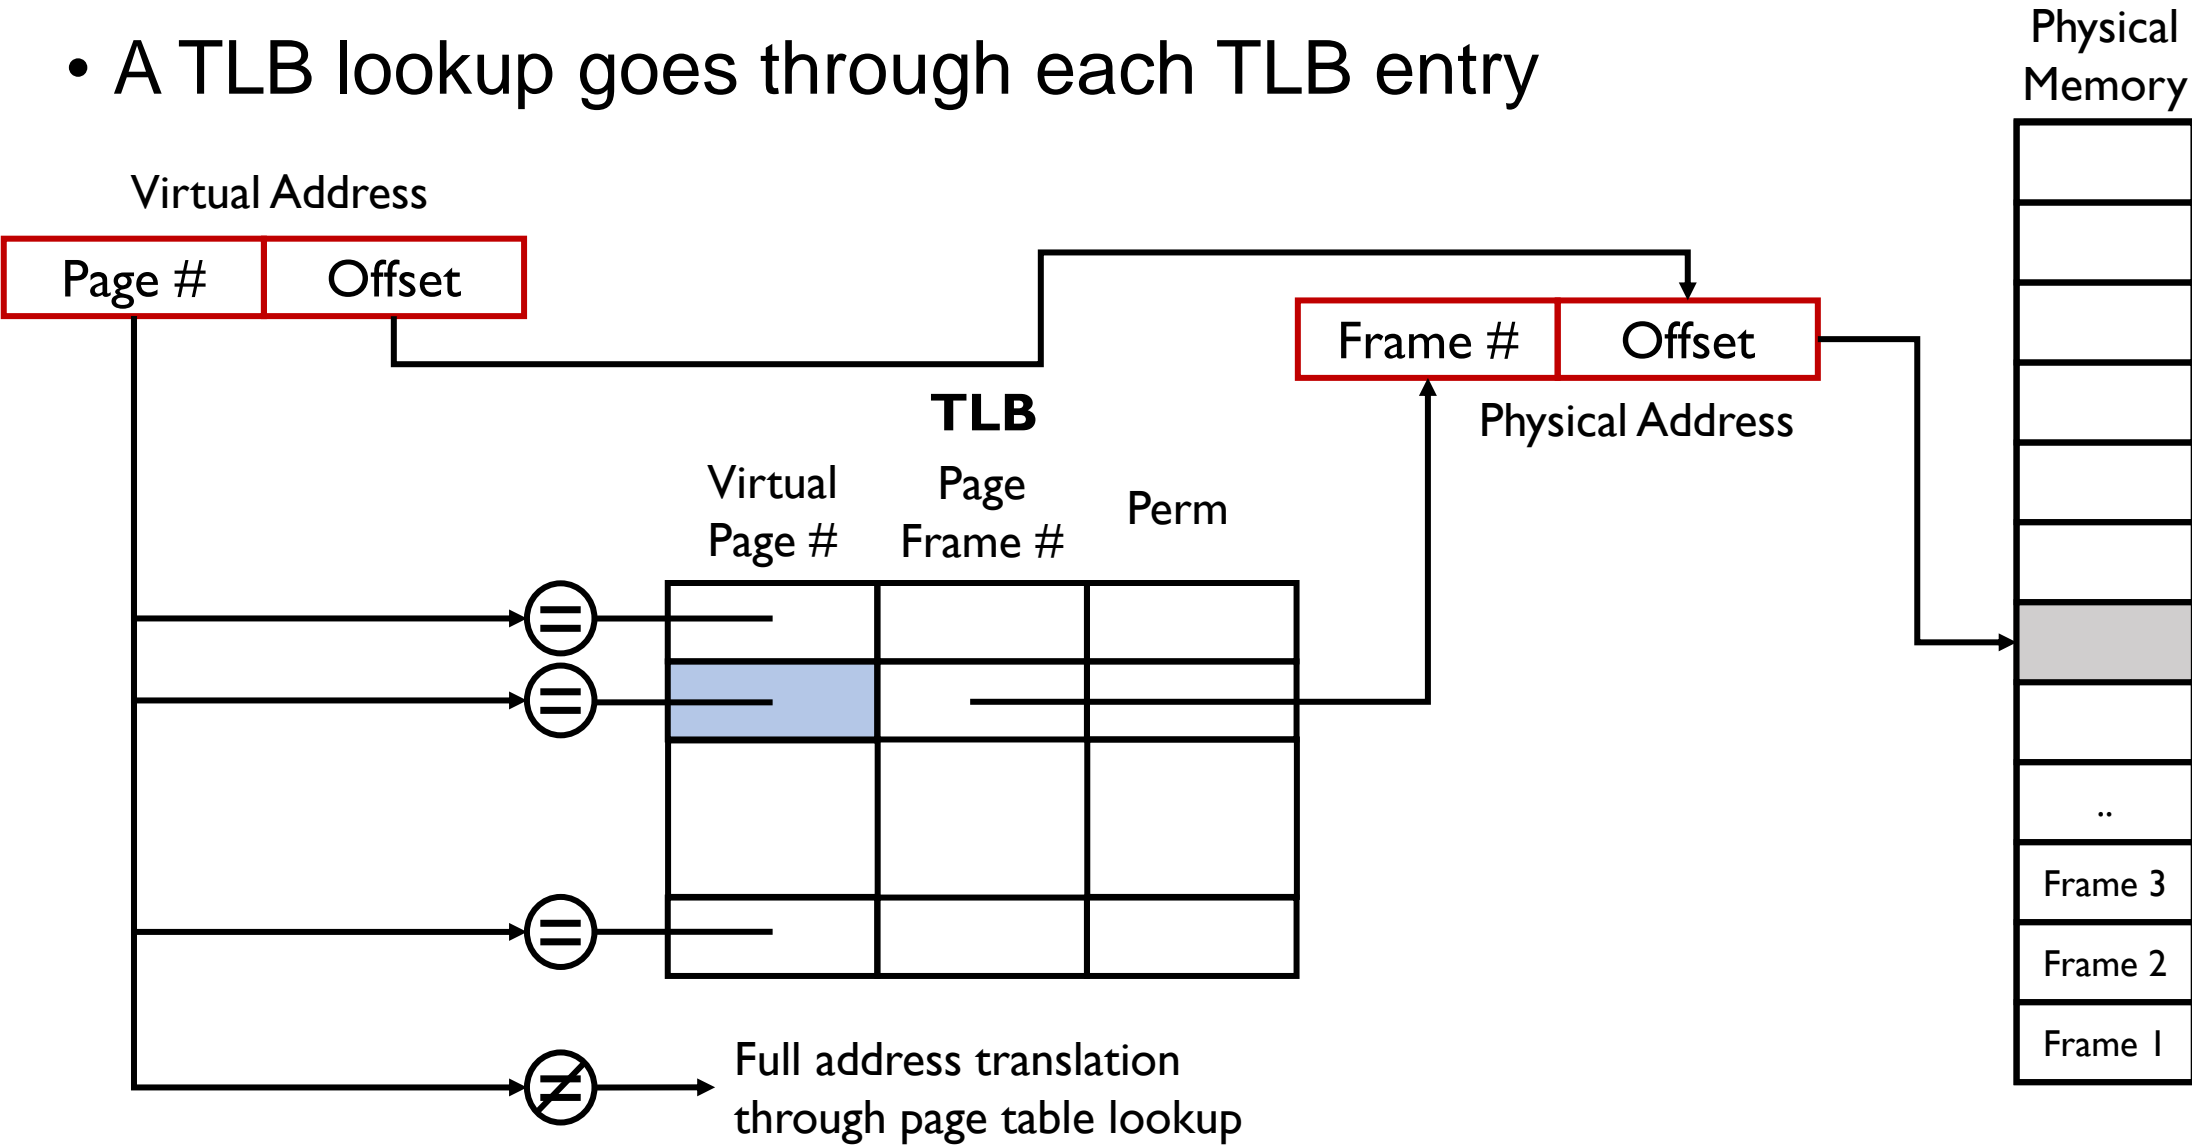
- Translation Lookaside Buffers (TLB, 转换检测缓冲区): a special cache within MMU that accelerates address translation

- The time and spatial locality. Who are they?
- Memory mapping is page-aligned.



Recap: TLB Lookup

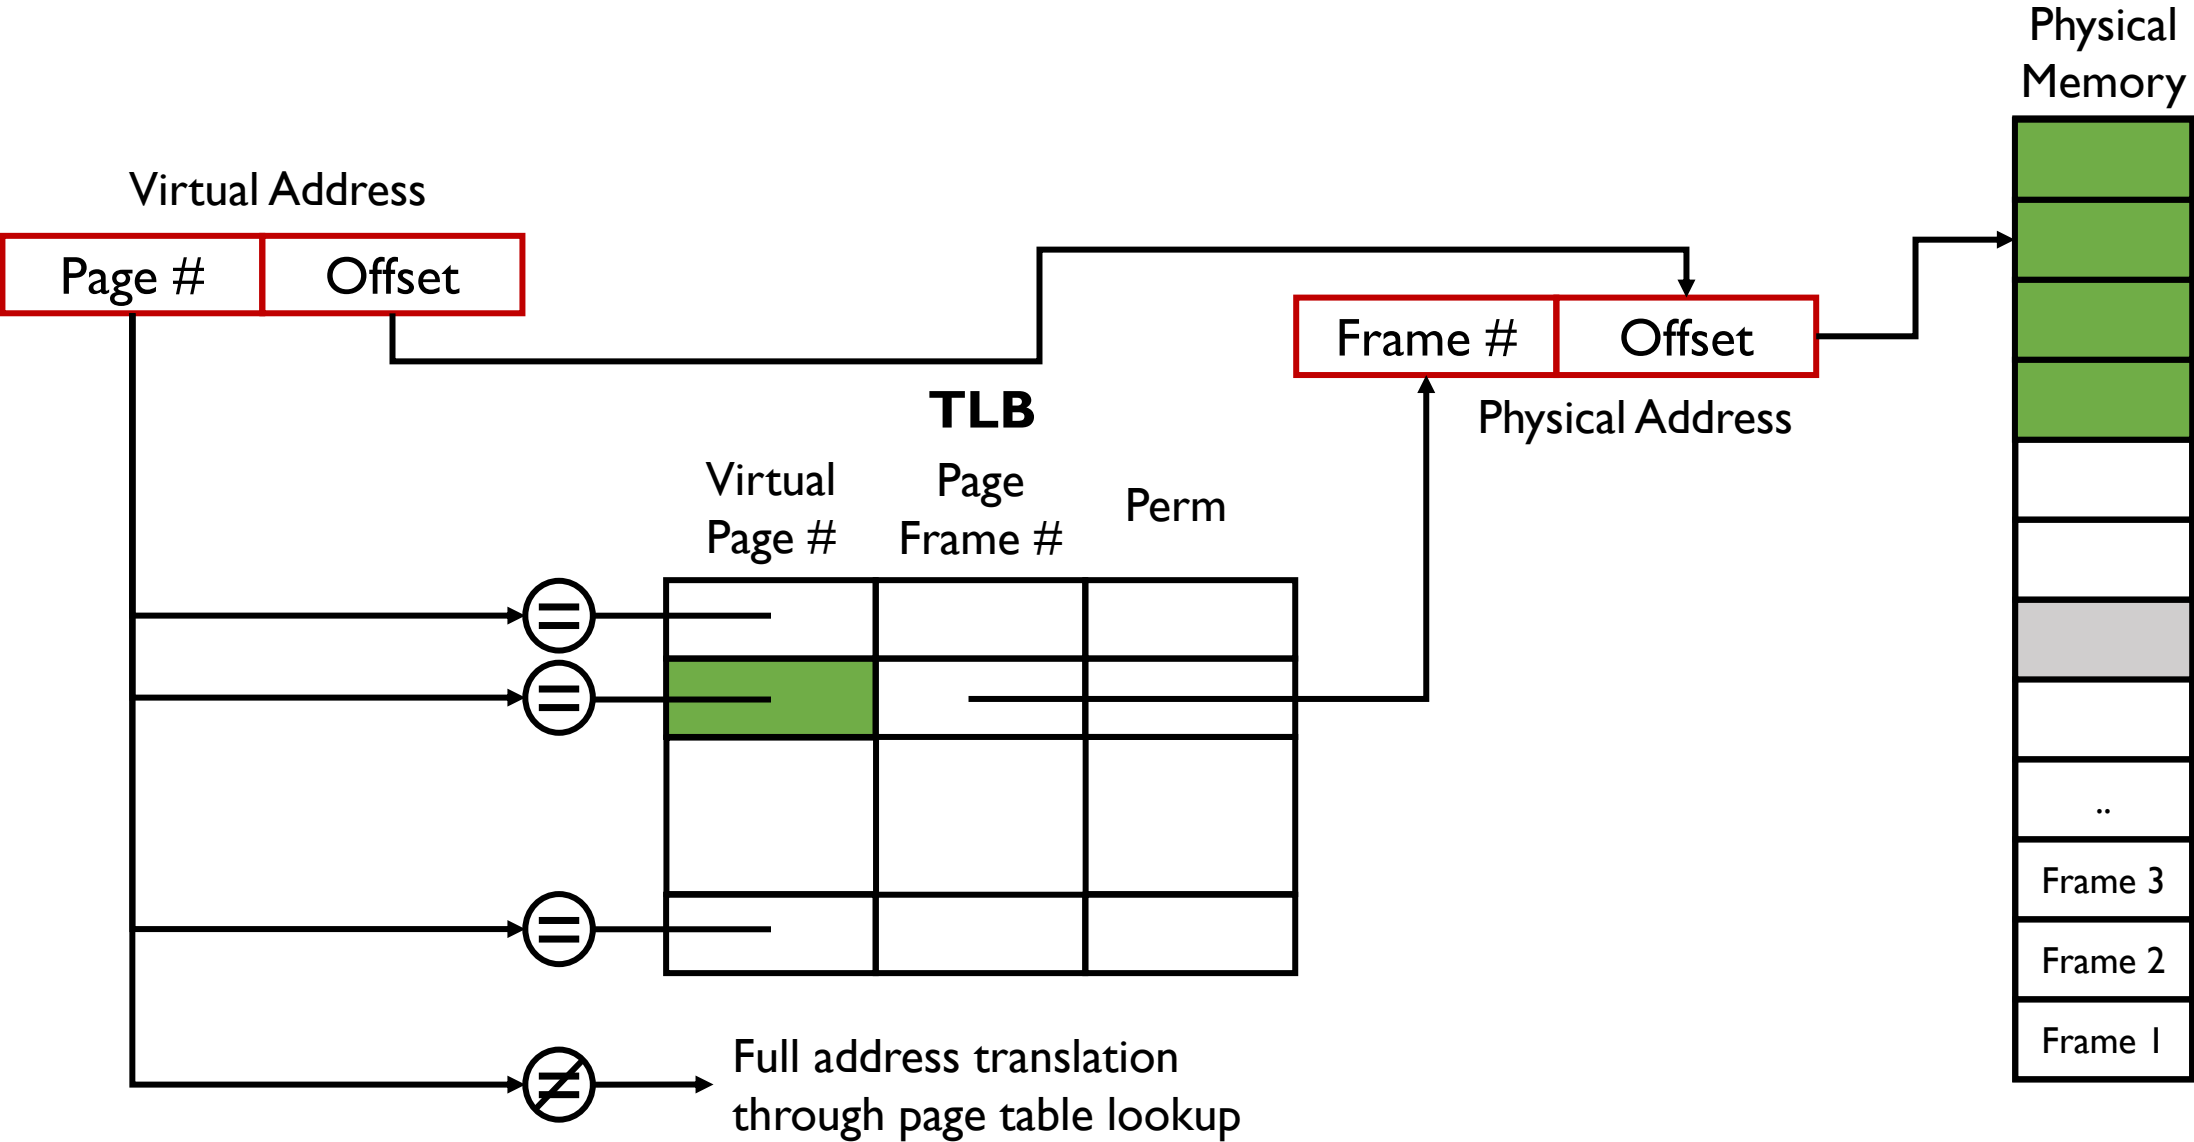
- A TLB lookup goes through each TLB entry



Recap: TLB Miss

- (Mostly) Hardware traversed page tables:
 - On TLB miss, hardware in MMU looks at current page table to fill TLB (may walk multiple levels)
 - ☐ If PTE valid, hardware fills TLB and processor never knows
 - ☐ If PTE marked as invalid, causes Page Fault, after which kernel decides what to do afterwards
- Software traversed Page tables (like MIPS)
 - On TLB miss, processor receives TLB fault
 - Kernel traverses page table to find PTE
 - ☐ If PTE valid, fills TLB and returns from fault
 - ☐ If PTE marked as invalid, internally calls Page Fault handler

Recap: Superpage




Recap: TLB Consistency

- Consistency (一致性) is a common issue for each cache: the cache must be always the same as the original data whenever the entries are modified.
 - Process context switch
 - Permission reduction
 - TLB shutdown

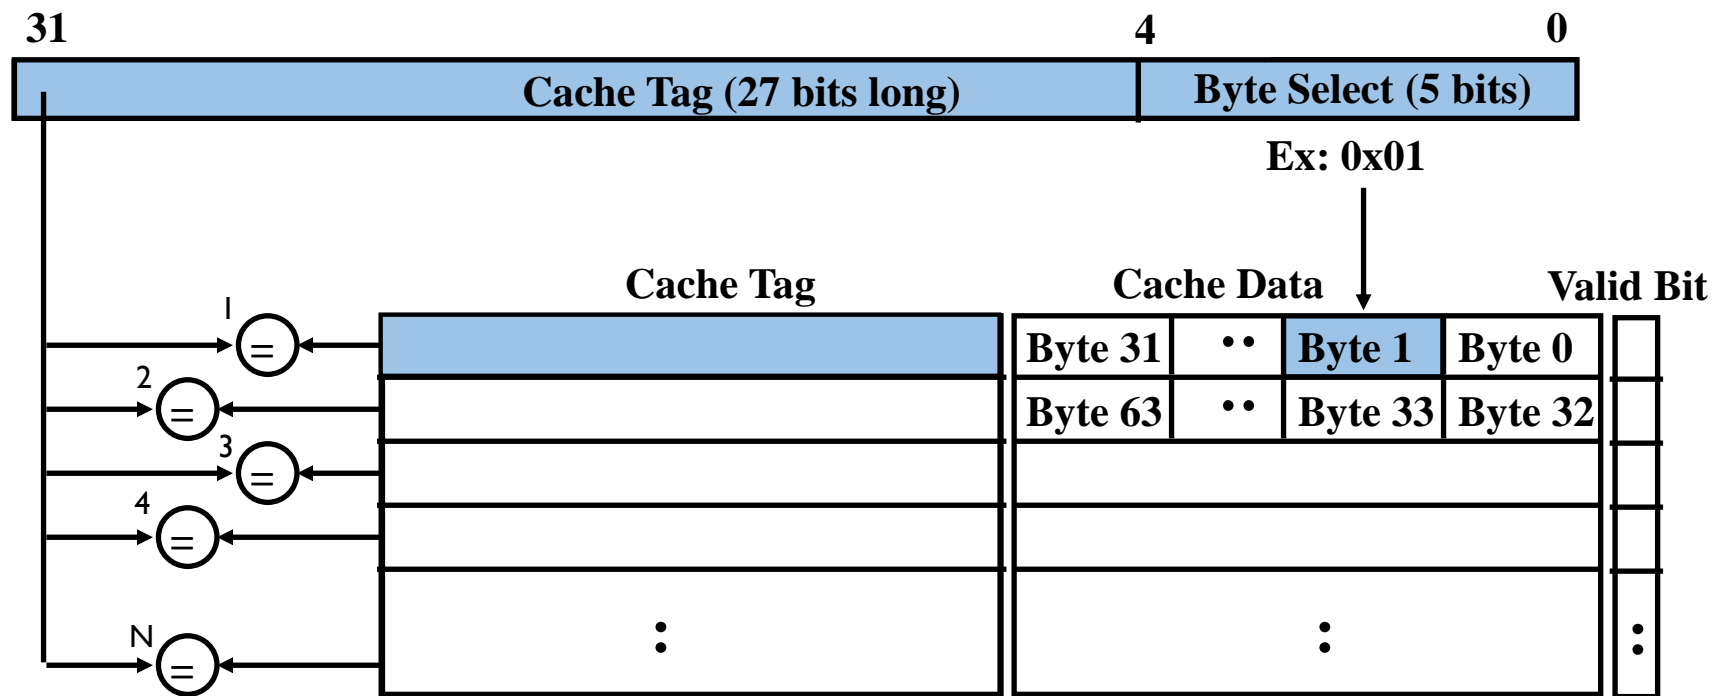
Recap: Cache Lookup

- Fully associative (全关联、完全关联): each address can be stored anywhere in the cache table
 - Direct mapped (直接映射): each address can be stored in one location in the cache table
 - N-way set associative (N路组关联): each address can be stored in one of N cache sets
-
- Tradeoffs: lookup speed and cache hit rate



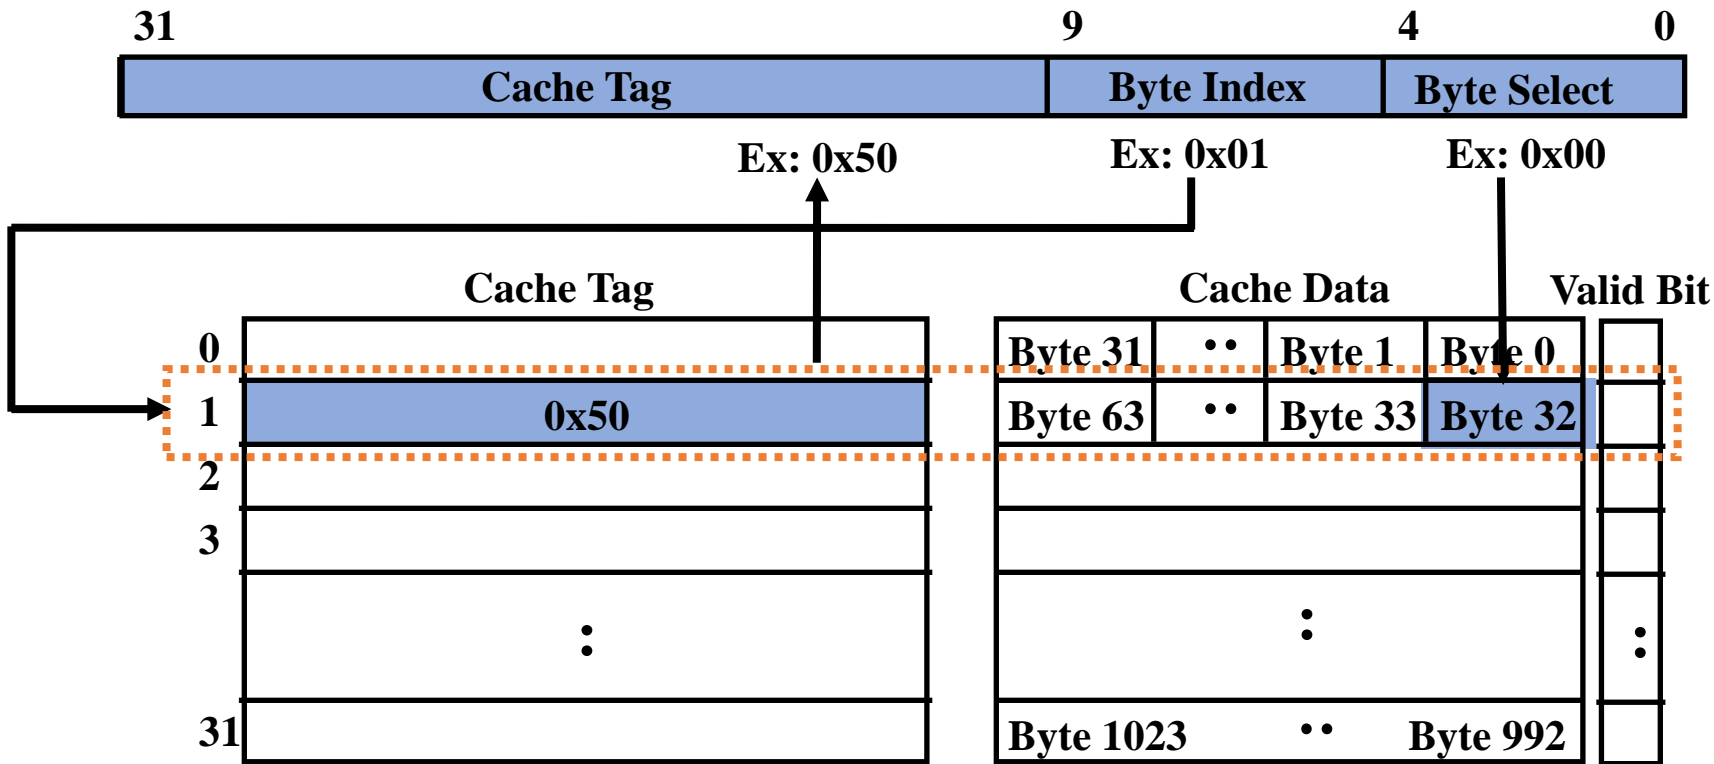
Recap: Fully Associative

- Compare the cache tag on each cache line
- Example: Block Size=32B blocks
 - We need Nx 27-bit comparators



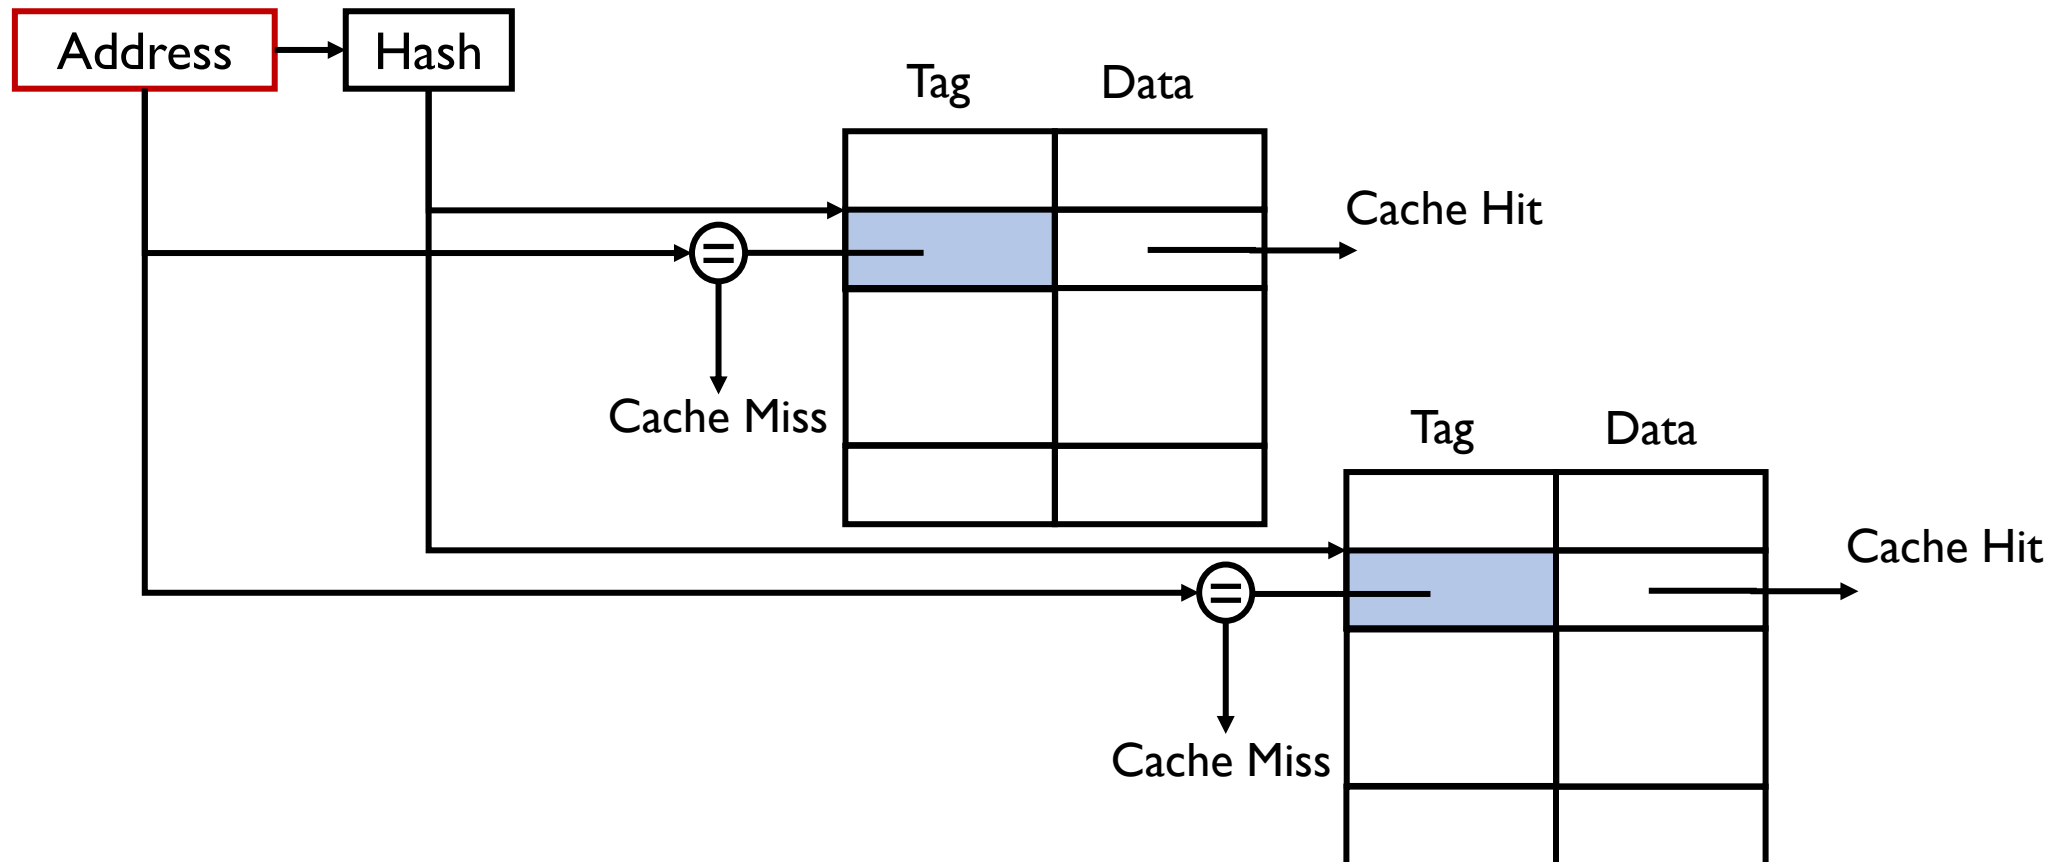
Recap: Direct Mapped

- Example: 1 KB Direct Mapped Cache with 32B Blocks
 - Index chooses potential block
 - Tag checked to verify block
 - Byte select chooses byte within block



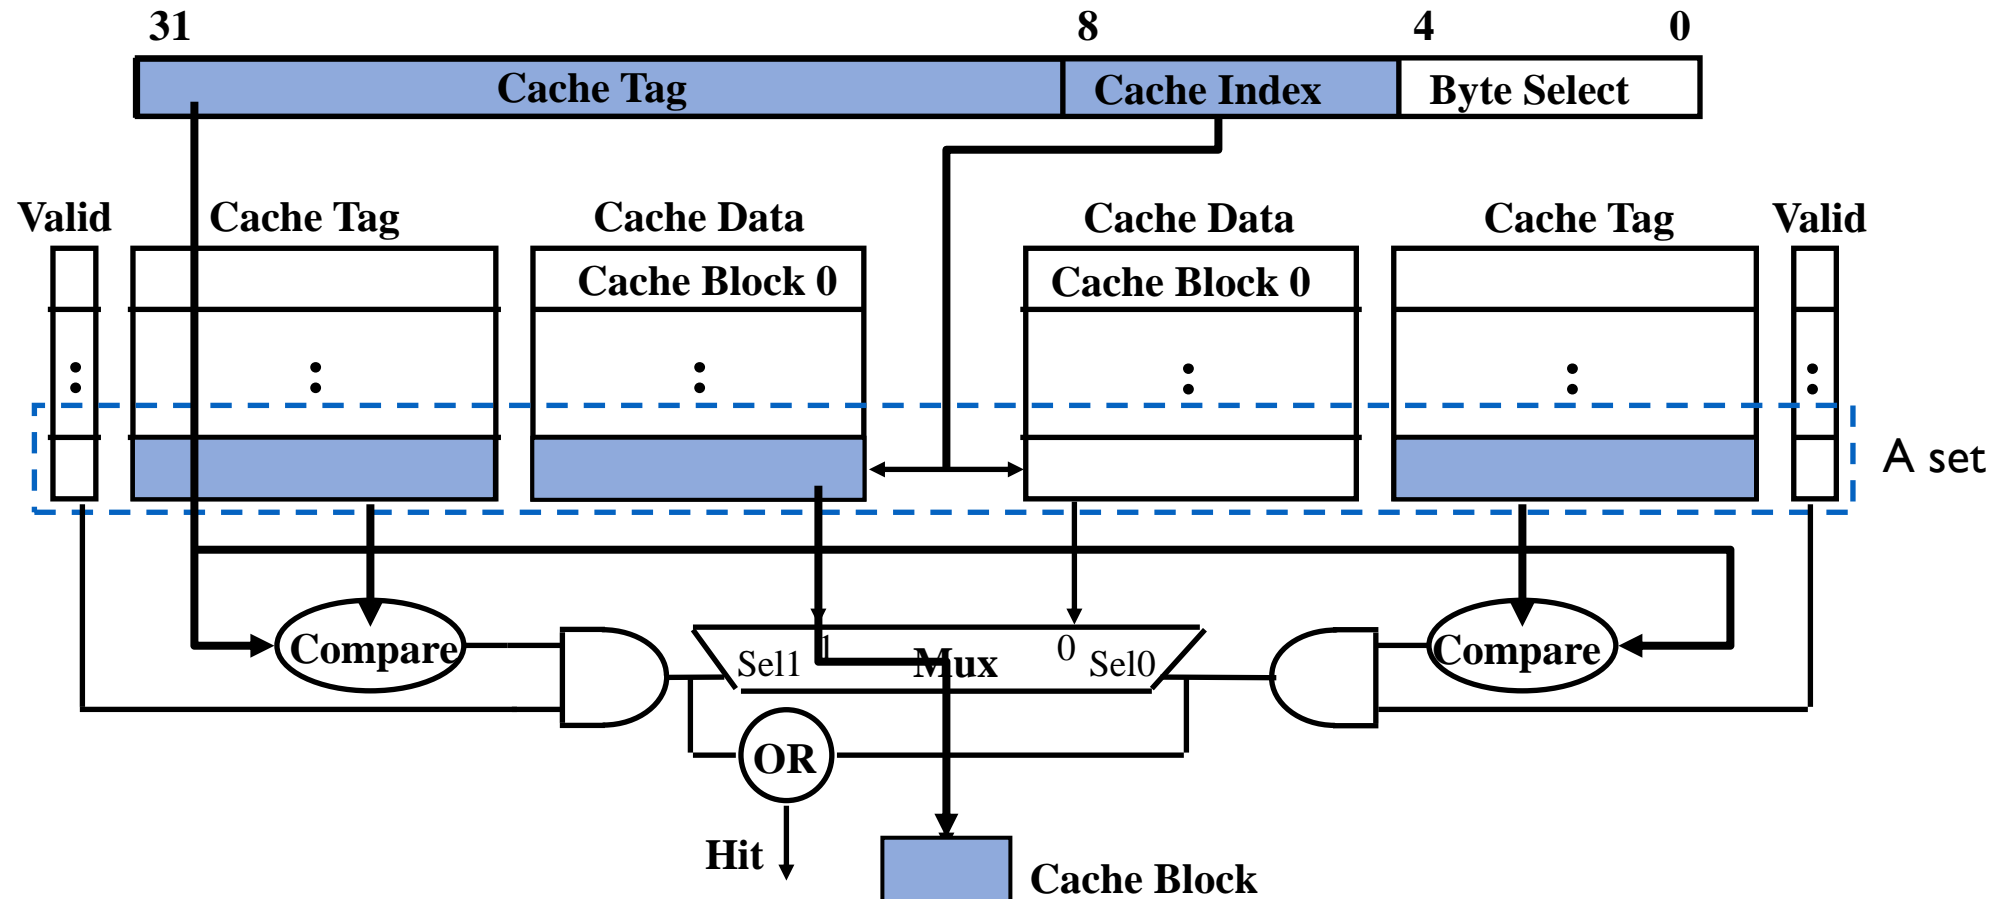
Recap: Set Associative

- N-way Set Associative: N entries per Cache Index
 - N direct mapped caches operates in parallel




Recap: Set Associative

- Example: two-way set associative cache
 - Cache Index selects a “set” from the cache
 - Two tags in the set are compared to input in parallel
 - Data is selected based on the tag result



Recap: Addressed Virtually or Physically?

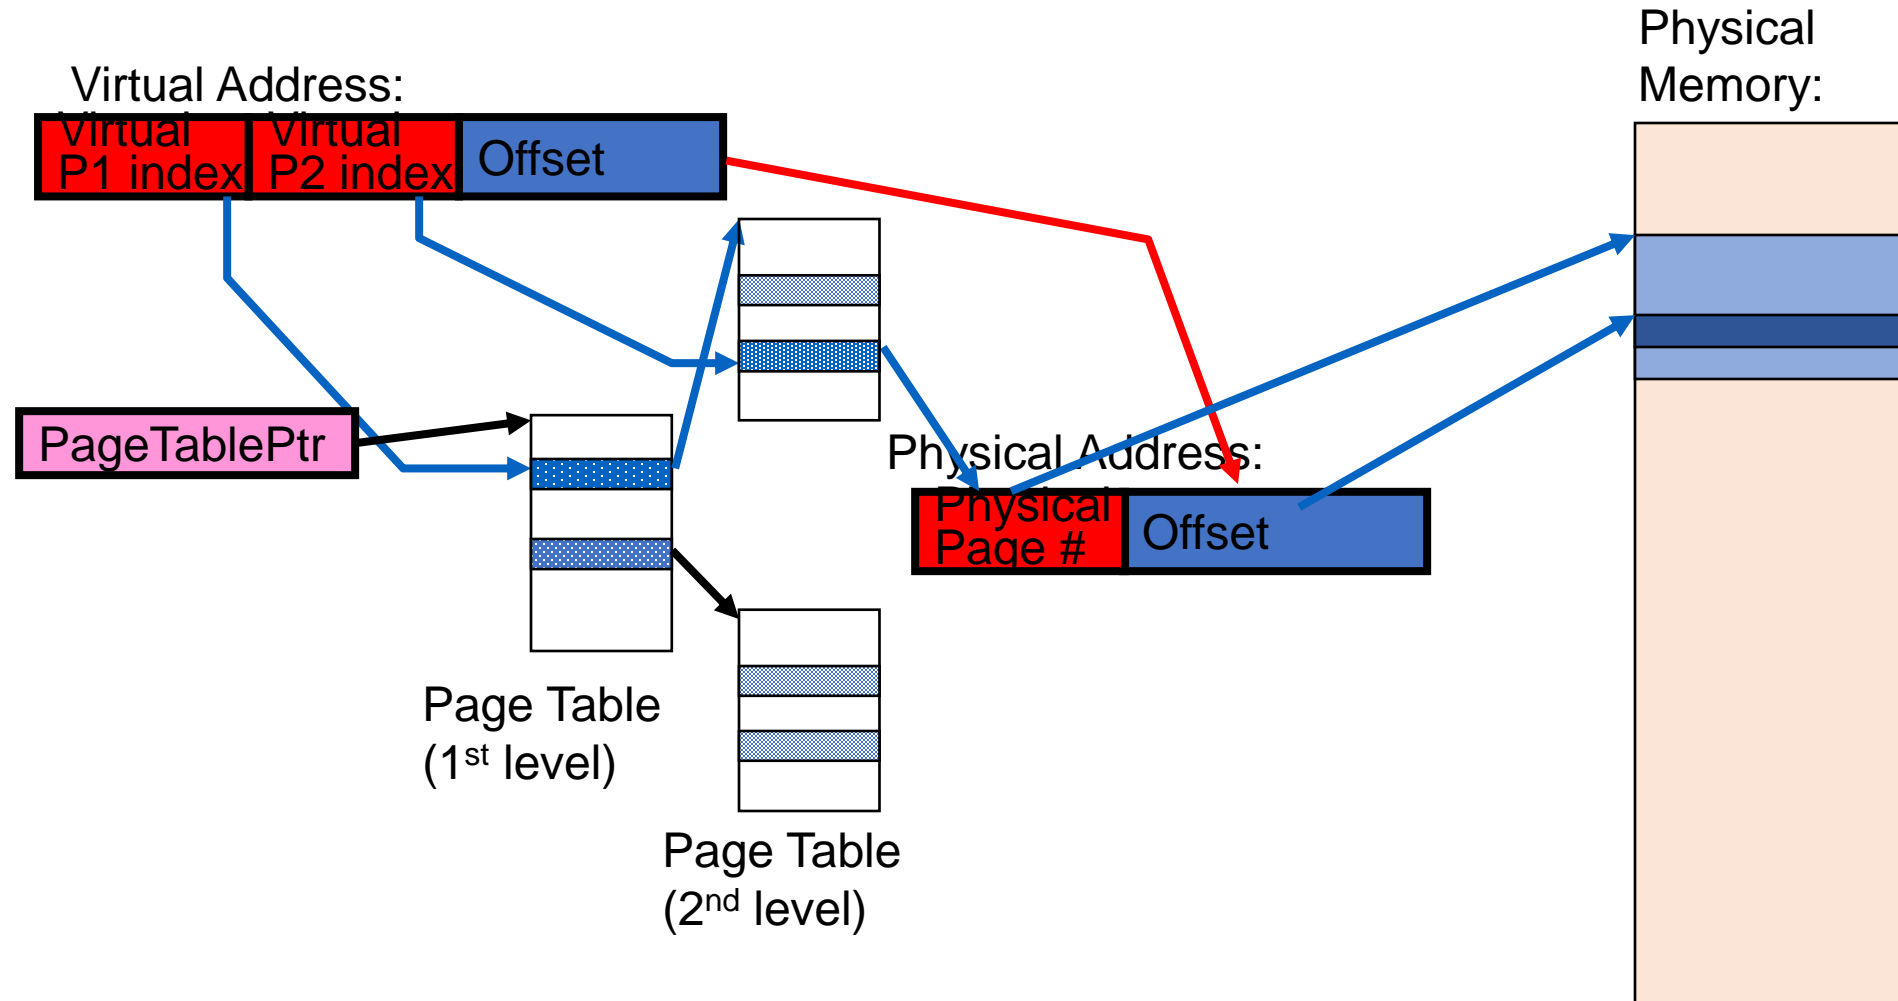
- The cache is addressed through virtual or physical address?
 - Note there are many levels of cache
- Every address access out of CPU is physical
 - The TLB miss cost is very high
 - Overlapping TLB and 1st-level cache as they are both in CPU



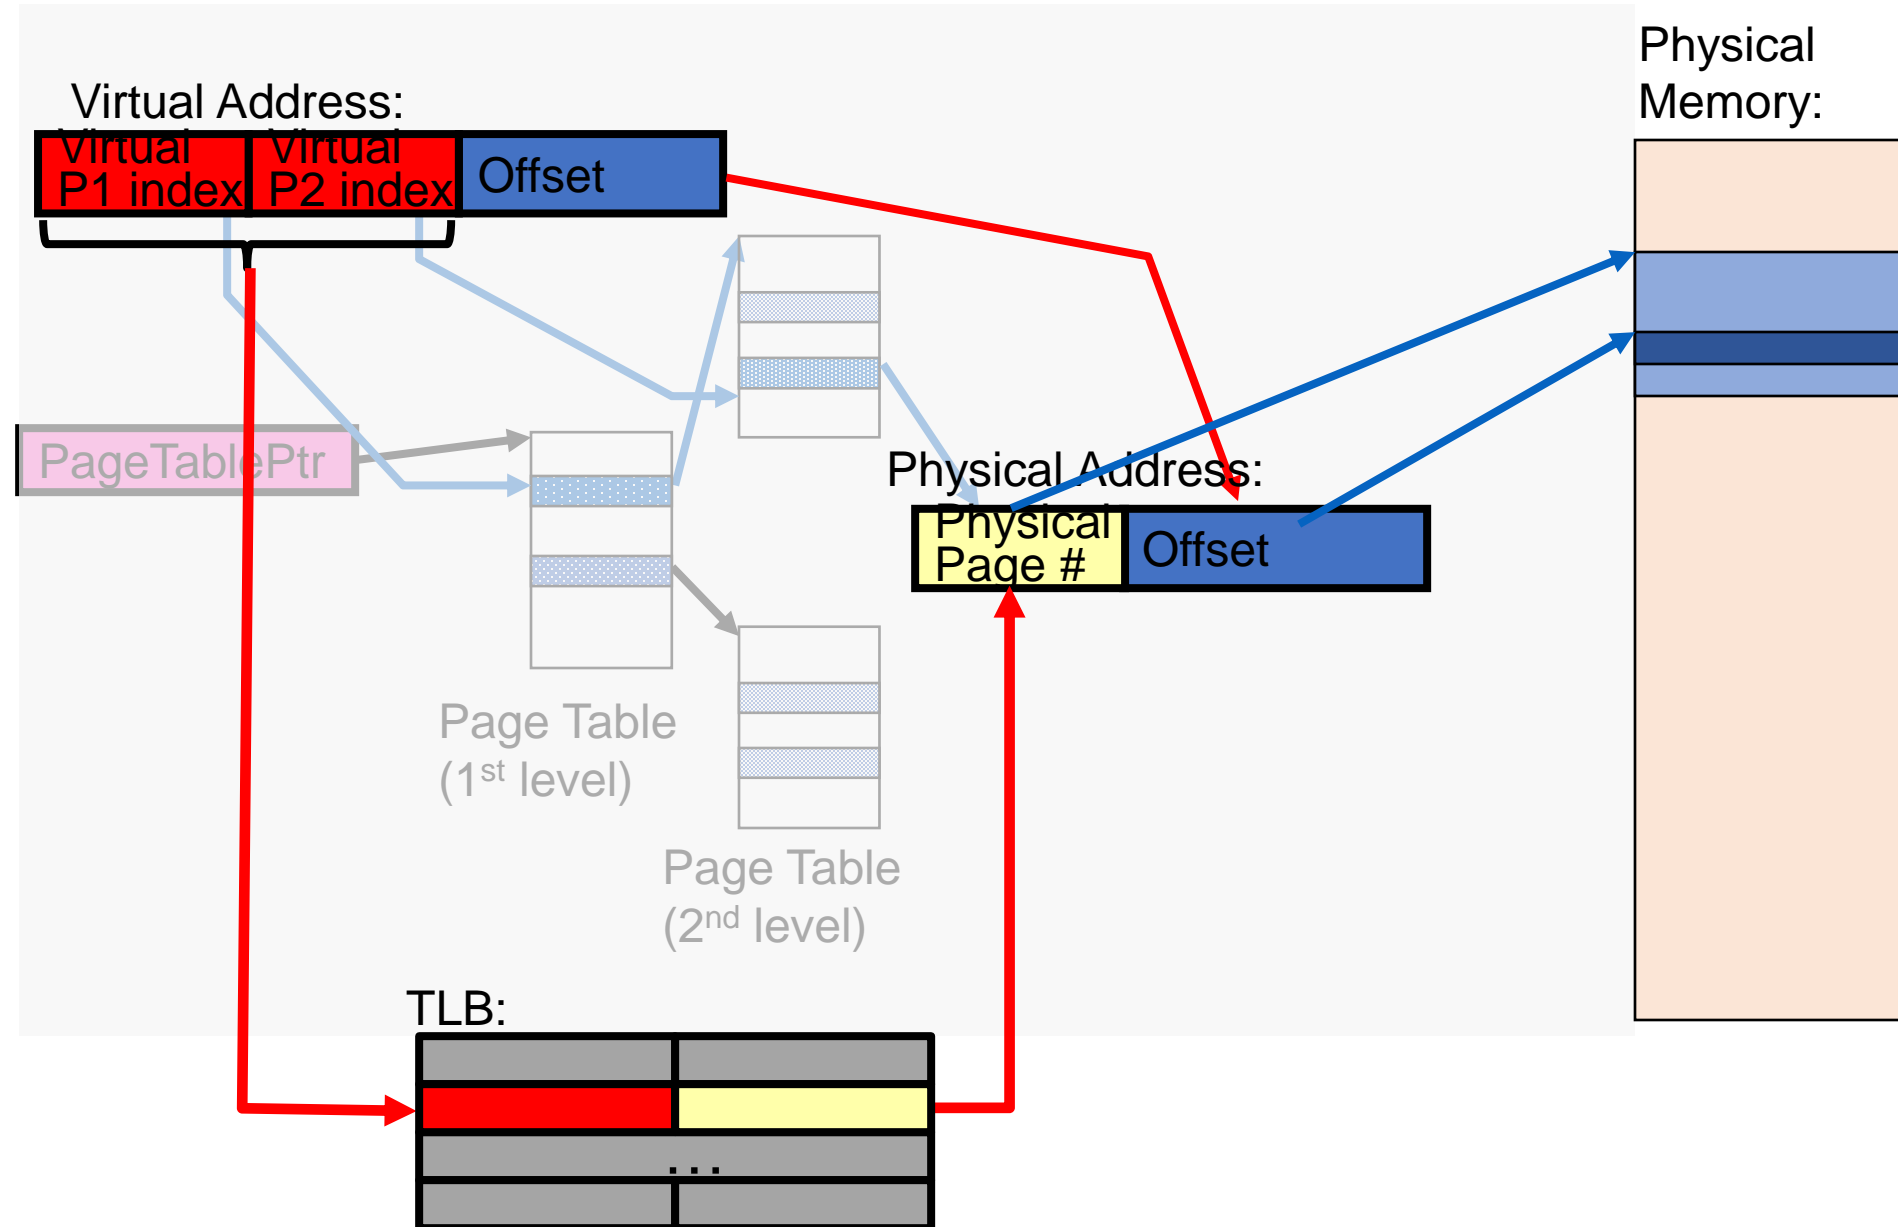
Recap: Overlapping TLB and Cache

- Key idea:
 - Offset in virtual address exactly covers the “cache index” and “byte select”
 - Thus can select the cached byte(s) in parallel to perform address translation
 - “Virtually indexed, physically tagged” (VIPT)
- Another option: virtually indexed, virtually tagged (VIVT)
 - Tags in cache are virtual addresses
 - Translation only happens on cache misses
 - What’s the problems?
- L1 is mostly VIPT, L2/L3 are mostly PIPT

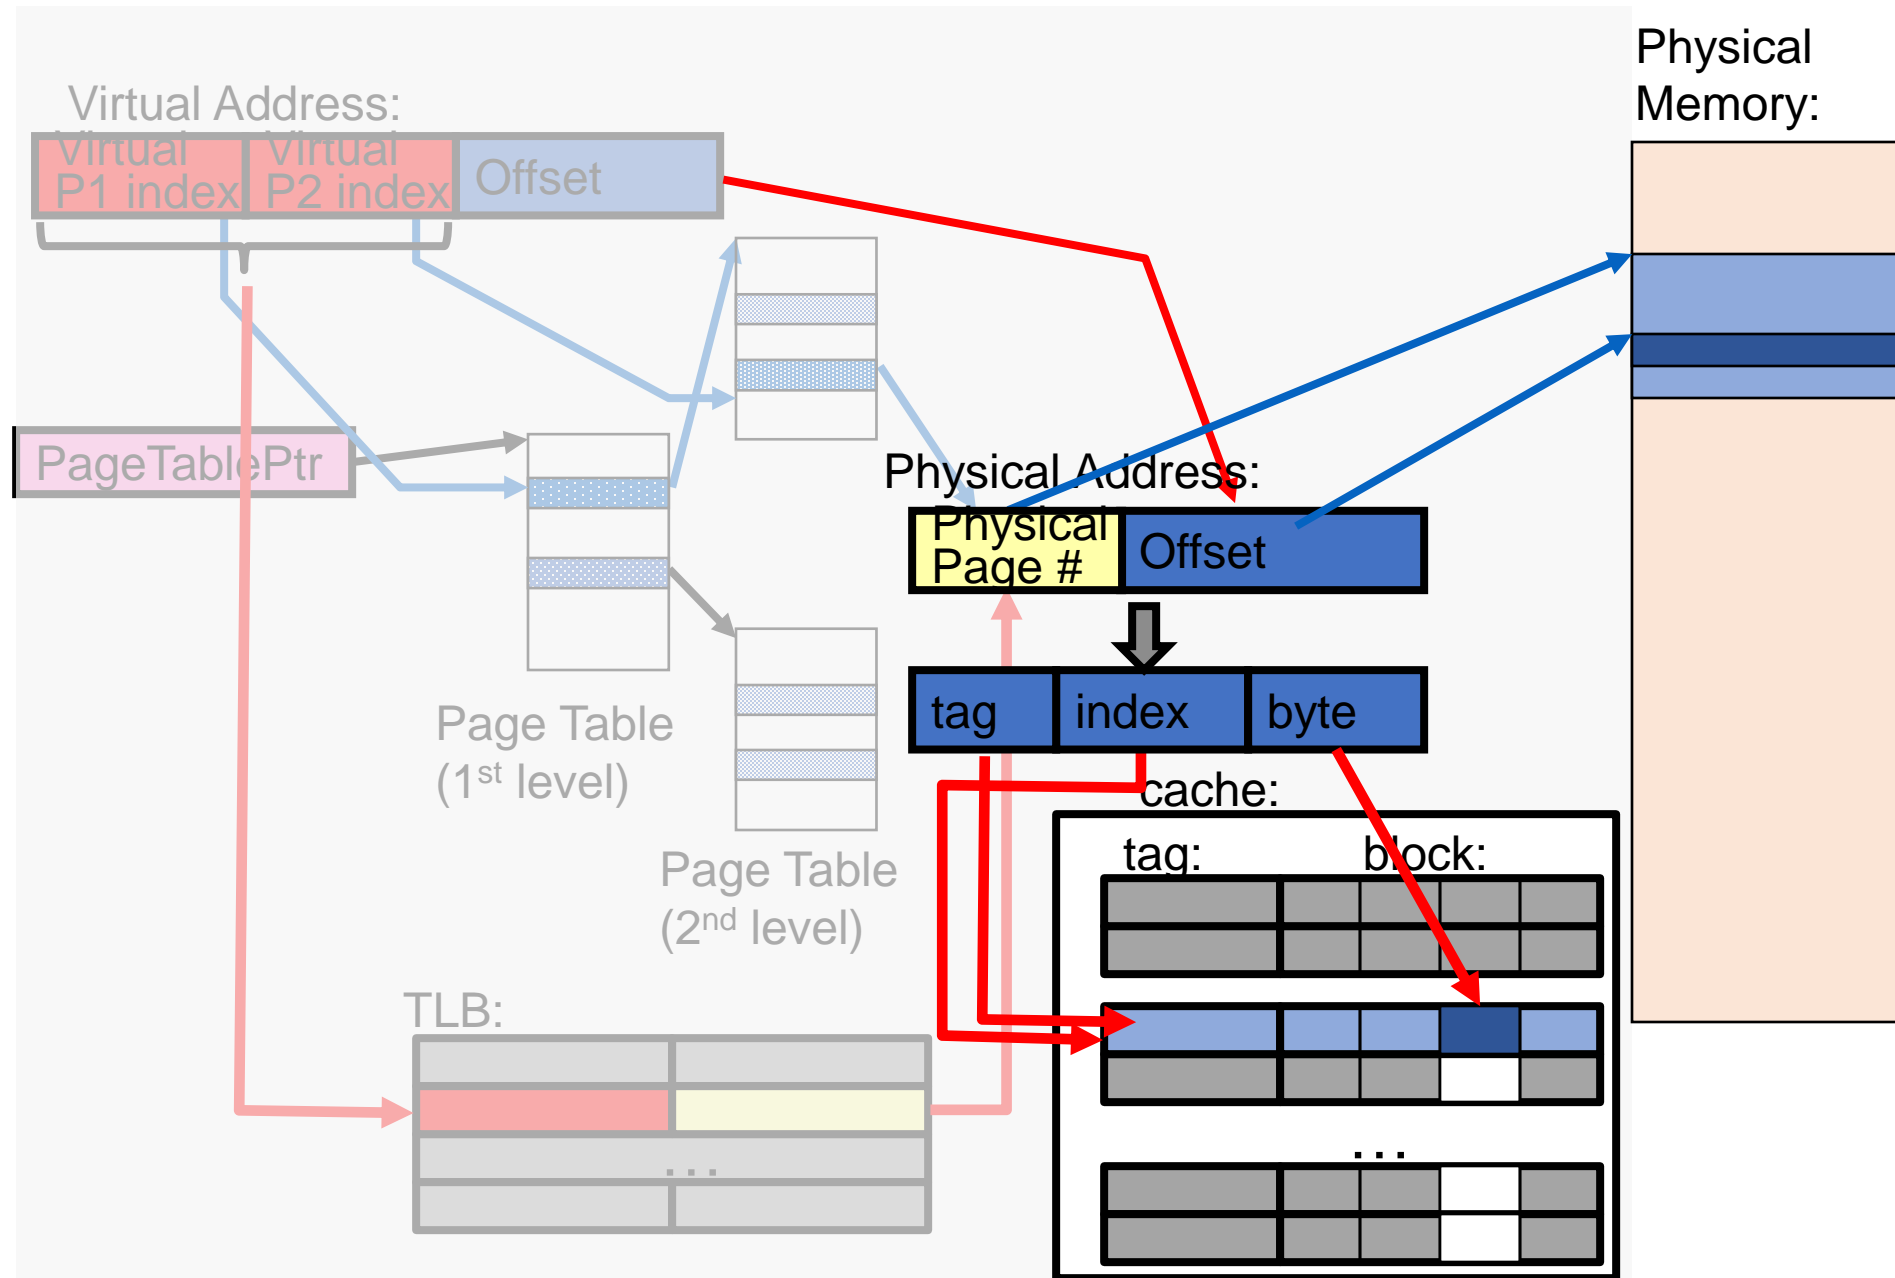
Recap: Putting Everything Together: Address Translation



Recap: Putting Everything Together: TLB

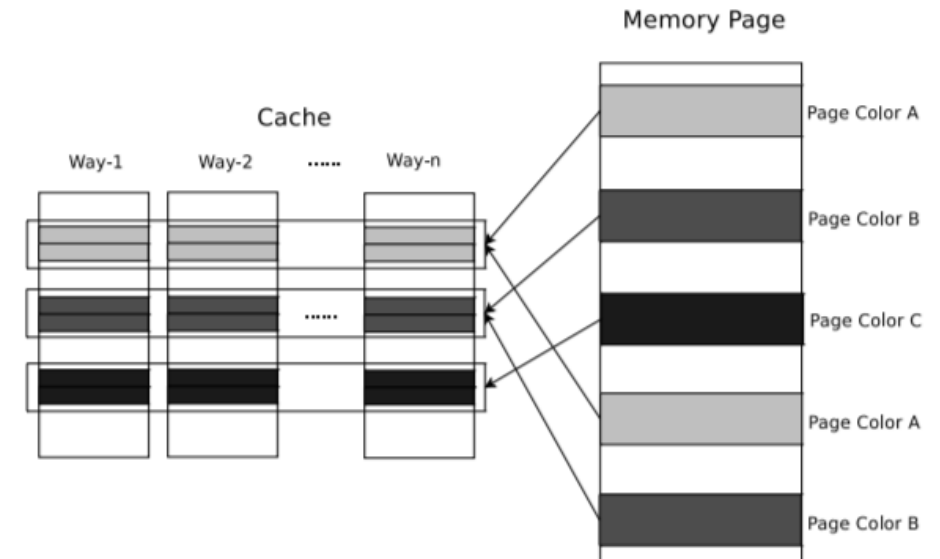



Recap: Putting Everything Together: Cache



Recap: Page Coloring

- Page Coloring or Cache Coloring (着色) technique helps reduce the cache miss in an app



Consider two consecutive pages used by an application:


- Their virtual set number must be different
- But their physical set number could be the same after translation (when the OS maps them to the physical pages whose page numbers have the same last 2 bits). In such a case, two addresses with the same offset within these two pages will in contention for the cache set.

Solutions

- Coloring the physical pages with the cache sets
- Maps the application pages to as many colors as possible (so less contention)

Cache Hierarchy

- Memory as cache for secondary disk



Demand Paging (需求分页)

- Modern programs require a lot of physical memory, but they don't use all their memory all of the time
 - 90-10 rule: programs spend 90% of their time in 10% of their code
 - Wasteful to require all of user's code to be in memory

Demand Paging (需求分页)

- Modern programs require a lot of physical memory, but they don't use all their memory all of the time
 - 90-10 rule: programs spend 90% of their time in 10% of their code
 - Wasteful to require all of user's code to be in memory
- Solution: use main memory as cache for disk
 - “lazy” memory allocation

Demand Paging (需求分页)

- Modern programs require a lot of physical memory, but they don't use all their memory all of the time
 - 90-10 rule: programs spend 90% of their time in 10% of their code
 - Wasteful to require all of user's code to be in memory
- Solution: use main memory as cache for disk
 - “lazy” memory allocation
- An illusion of infinite memory
 - In-use virtual memory can be bigger than physical memory
 - Combined memory of running processes much larger than physical memory
 - ❑ More programs fit into memory, allowing more concurrency
 - **Principle: page table for transparent management**

Demand Paging as Cache

- What is block size?
 - 1 page
- What is organization of this cache (i.e. direct-mapped, set-associative, fully-associative)?
 - Fully associative: arbitrary virtual \rightarrow physical mapping
- How do we find a page in the cache when look for it?
 - First check TLB, then page-table traversal
- What is page replacement policy? (i.e. LRU, Random...)
 - This requires more explanation... (kinda LRU)
- What happens on a miss?
 - Go to lower level to fill miss (i.e. disk)
- What happens on a write? (write-through, write back)
 - Write-back – need dirty bit!

Memory-mapped Files

- Memory-mapped Files (内存映射文件) is a segment of virtual memory that has been assigned a direct byte-for-byte correlation with some portion of a file or file-like resource
 - A special case of demand paging
 - A replacement for syscall `read()/write()`

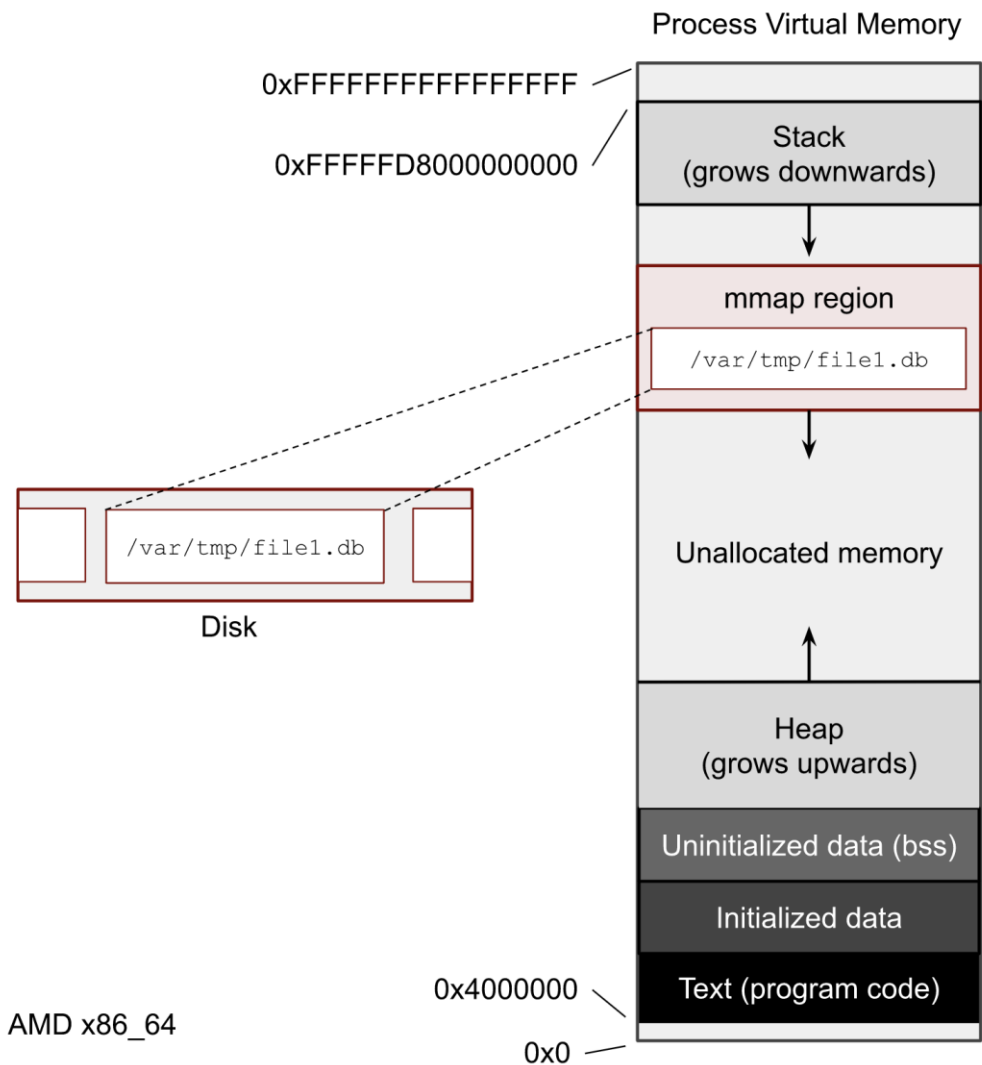
```
#include <sys/mman.h>

void *mmap(void *addr, size_t length, int prot, int flags,
           int fd, off_t offset);
int munmap(void *addr, size_t length);
```

`mmap()`: creates a new mapping in the virtual address space of the calling process. The virtual address starts at addr with length length. The contents of a file mapping are initialized using length bytes starting at offset offset in the file (or other object) referred to by the file descriptor fd.

- If `addr` is `NULL`, the OS picks a location
- Return value: the address of new mapping

Memory-mapped Files



```
int main() {
    int fd;
    char *mapped_data;
    struct stat file_stat;

    // Open the file for reading and writing
    fd = open("example.txt", O_RDWR);

    // Get file size
    if (fstat(fd, &file_stat) < 0) {
        return -1;
    }

    // Map the file into memory
    mapped_data = mmap(NULL, file_stat.st_size, PROT_READ |
PROT_WRITE, MAP_SHARED, fd, 0);

    // Modify the file in memory
    strncpy(mapped_data, "Modified", 8);

    // Sync changes to disk
    if (msync(mapped_data, file_stat.st_size, MS_SYNC) == -1) {
        return -1;
    }

    // Unmap the file and close fd
    if (munmap(mapped_data, file_stat.st_size) == -1) {
        return -1;
    }
    close(fd);
    return 0;
}
```

Memory-mapped Files

- PROS

- Transparency – the program can use pointers to access those data
- Zero copy I/O – the OS just changes the page table entries without copying the data into memory; read()/write() needs to copy the data twice (disk-kernel-user)
- Pipelining – the program can start executing as soon as the page table has been set
- Interprocess communication – sharing becomes easy
- Large files – which pages shall be in memory? OS handles it for you

- CONS

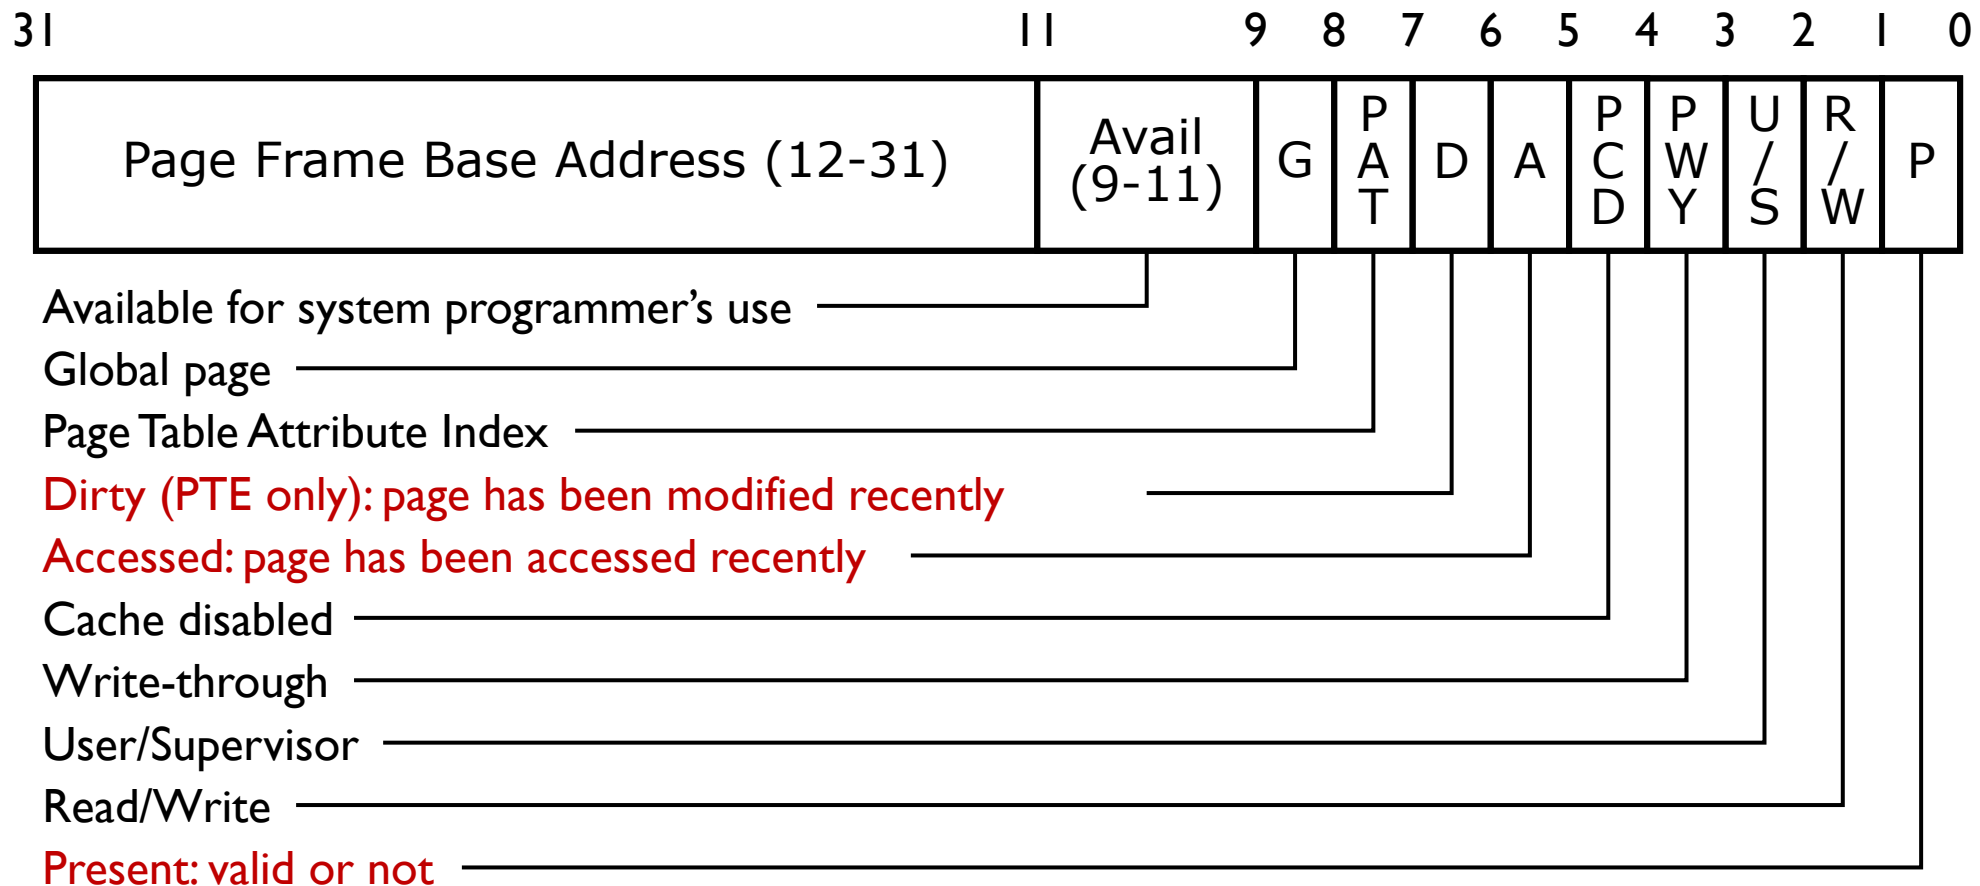
- Frequent page faults
- A few more..

When to use mmap

- Use mmap when:
 - Random Access: access data in a non-sequential manner
 - Large Files: for very large files that may not fit into memory
 - Multiple Processes: data sharing across processes
 - Memory-Mapped I/O and automatic caching
- Instead, use read/write when:
 - Small files
 - Streaming data access
 - Portability: not every OS has mmap!

Implementation of Memory-mapped Files

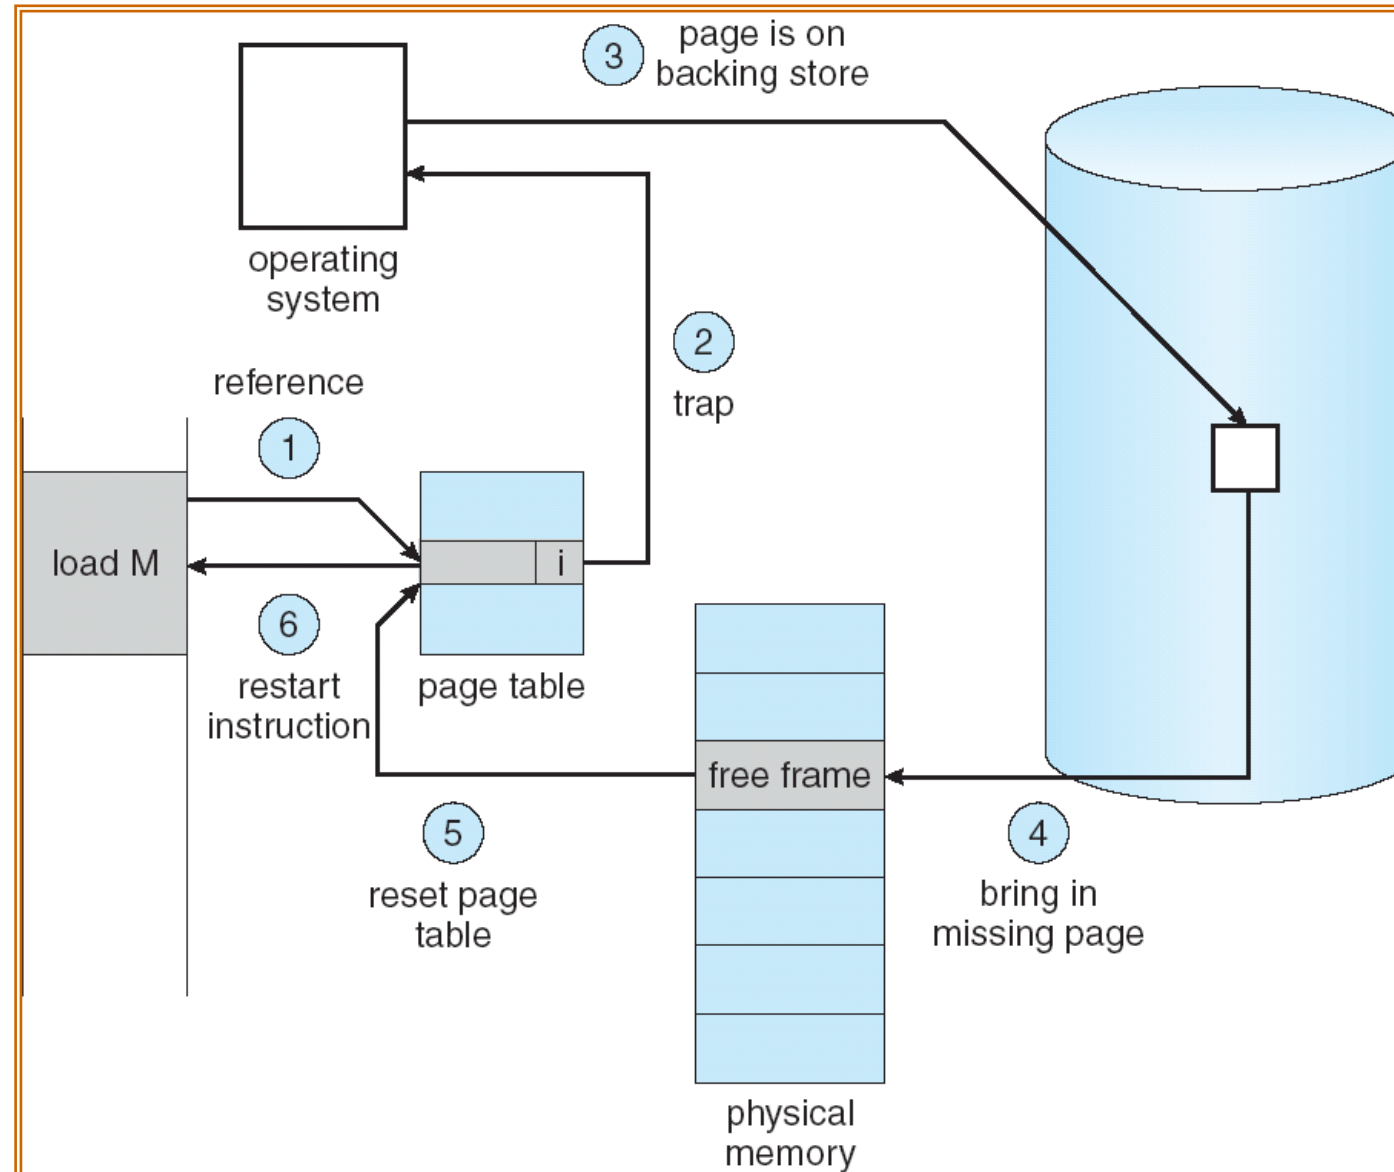
- Set up mapping
 - Initialize the page table entries and setting them to invalid



Implementation of Memory-mapped Files

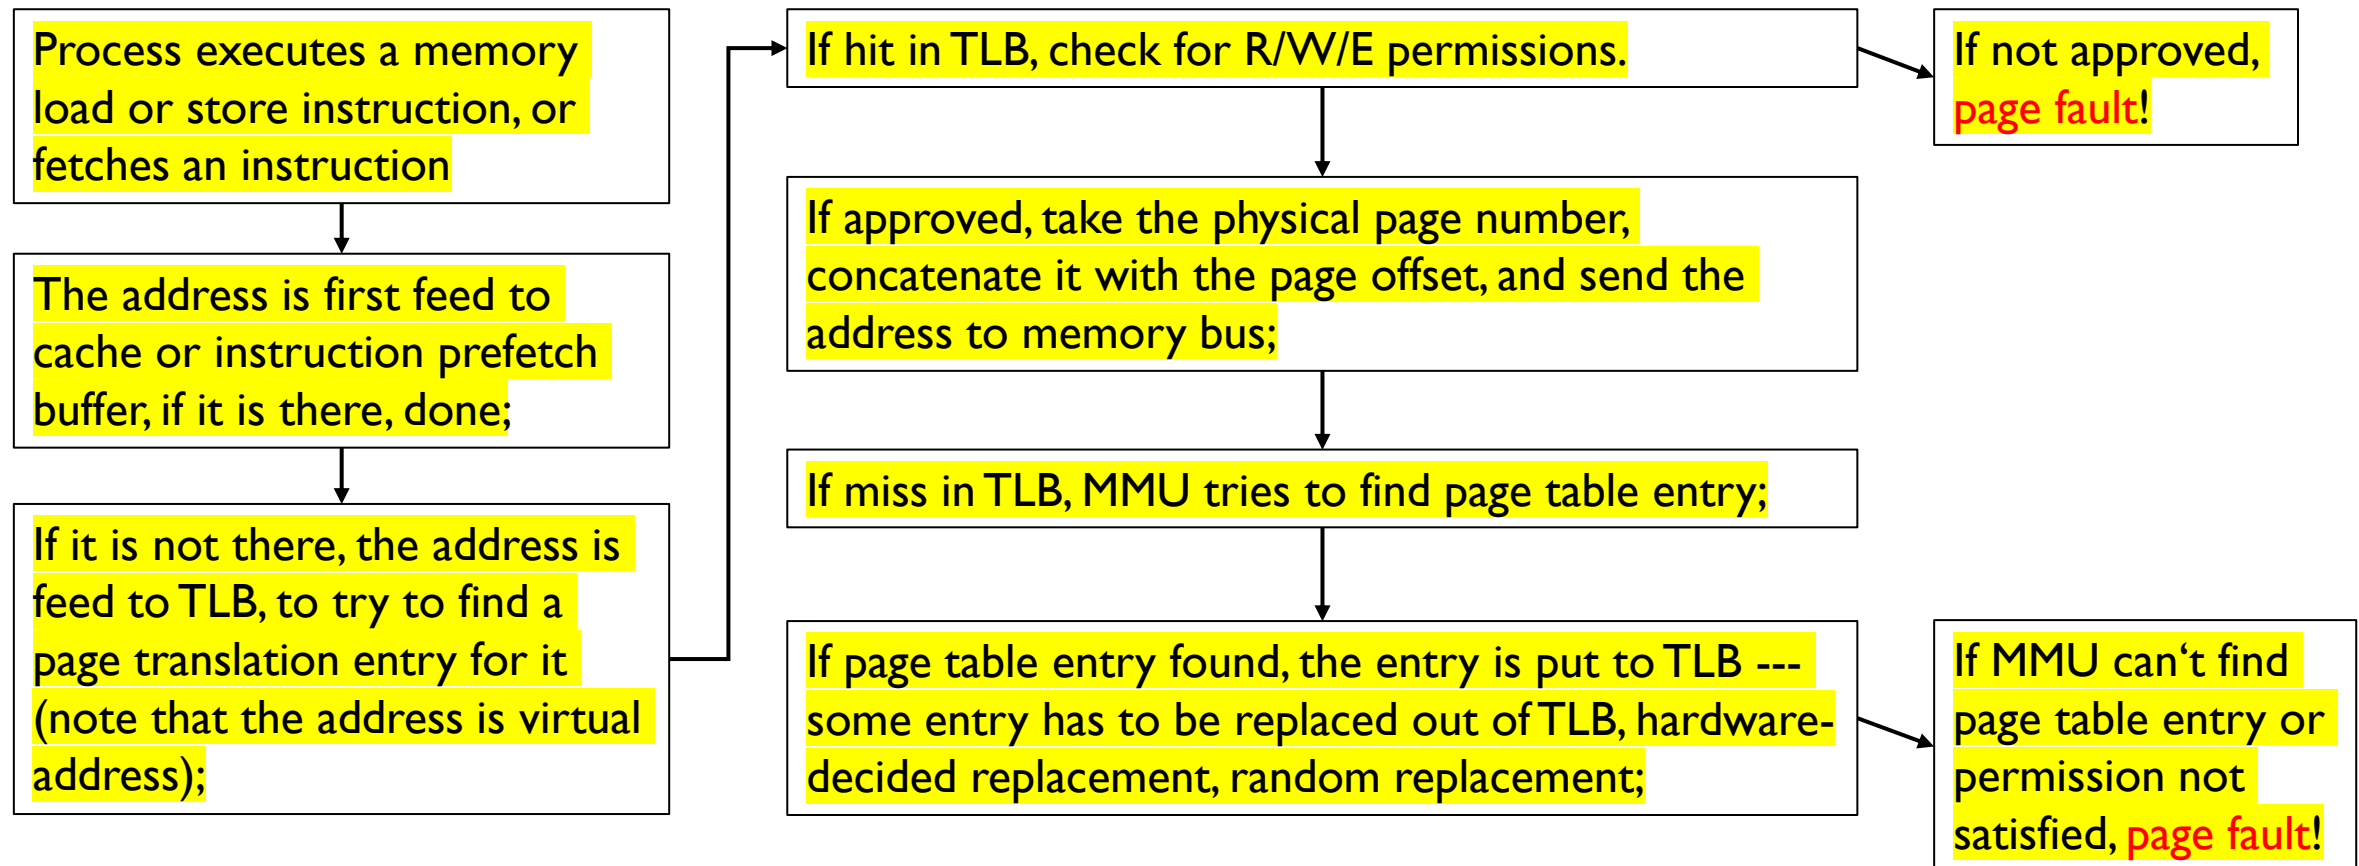
- When program accesses an invalid address
 1. [MMU] TLB miss; full page table lookup
 2. [MMU + OS] Trapping into page fault handler
 3. [OS] Convert virtual address to file offset
 4. [OS] Allocate a new page frame in memory
 5. [OS] Read data from disk to the memory (blocked)
 6. [CPU] Disk interrupt when read completes
 7. [OS] Updating page table by marking the entry as valid
 8. [OS] Resume process
 9. [MMU] TLB miss; full page table lookup
 10. [MMU] TLB update

Implementation of Memory-mapped Files



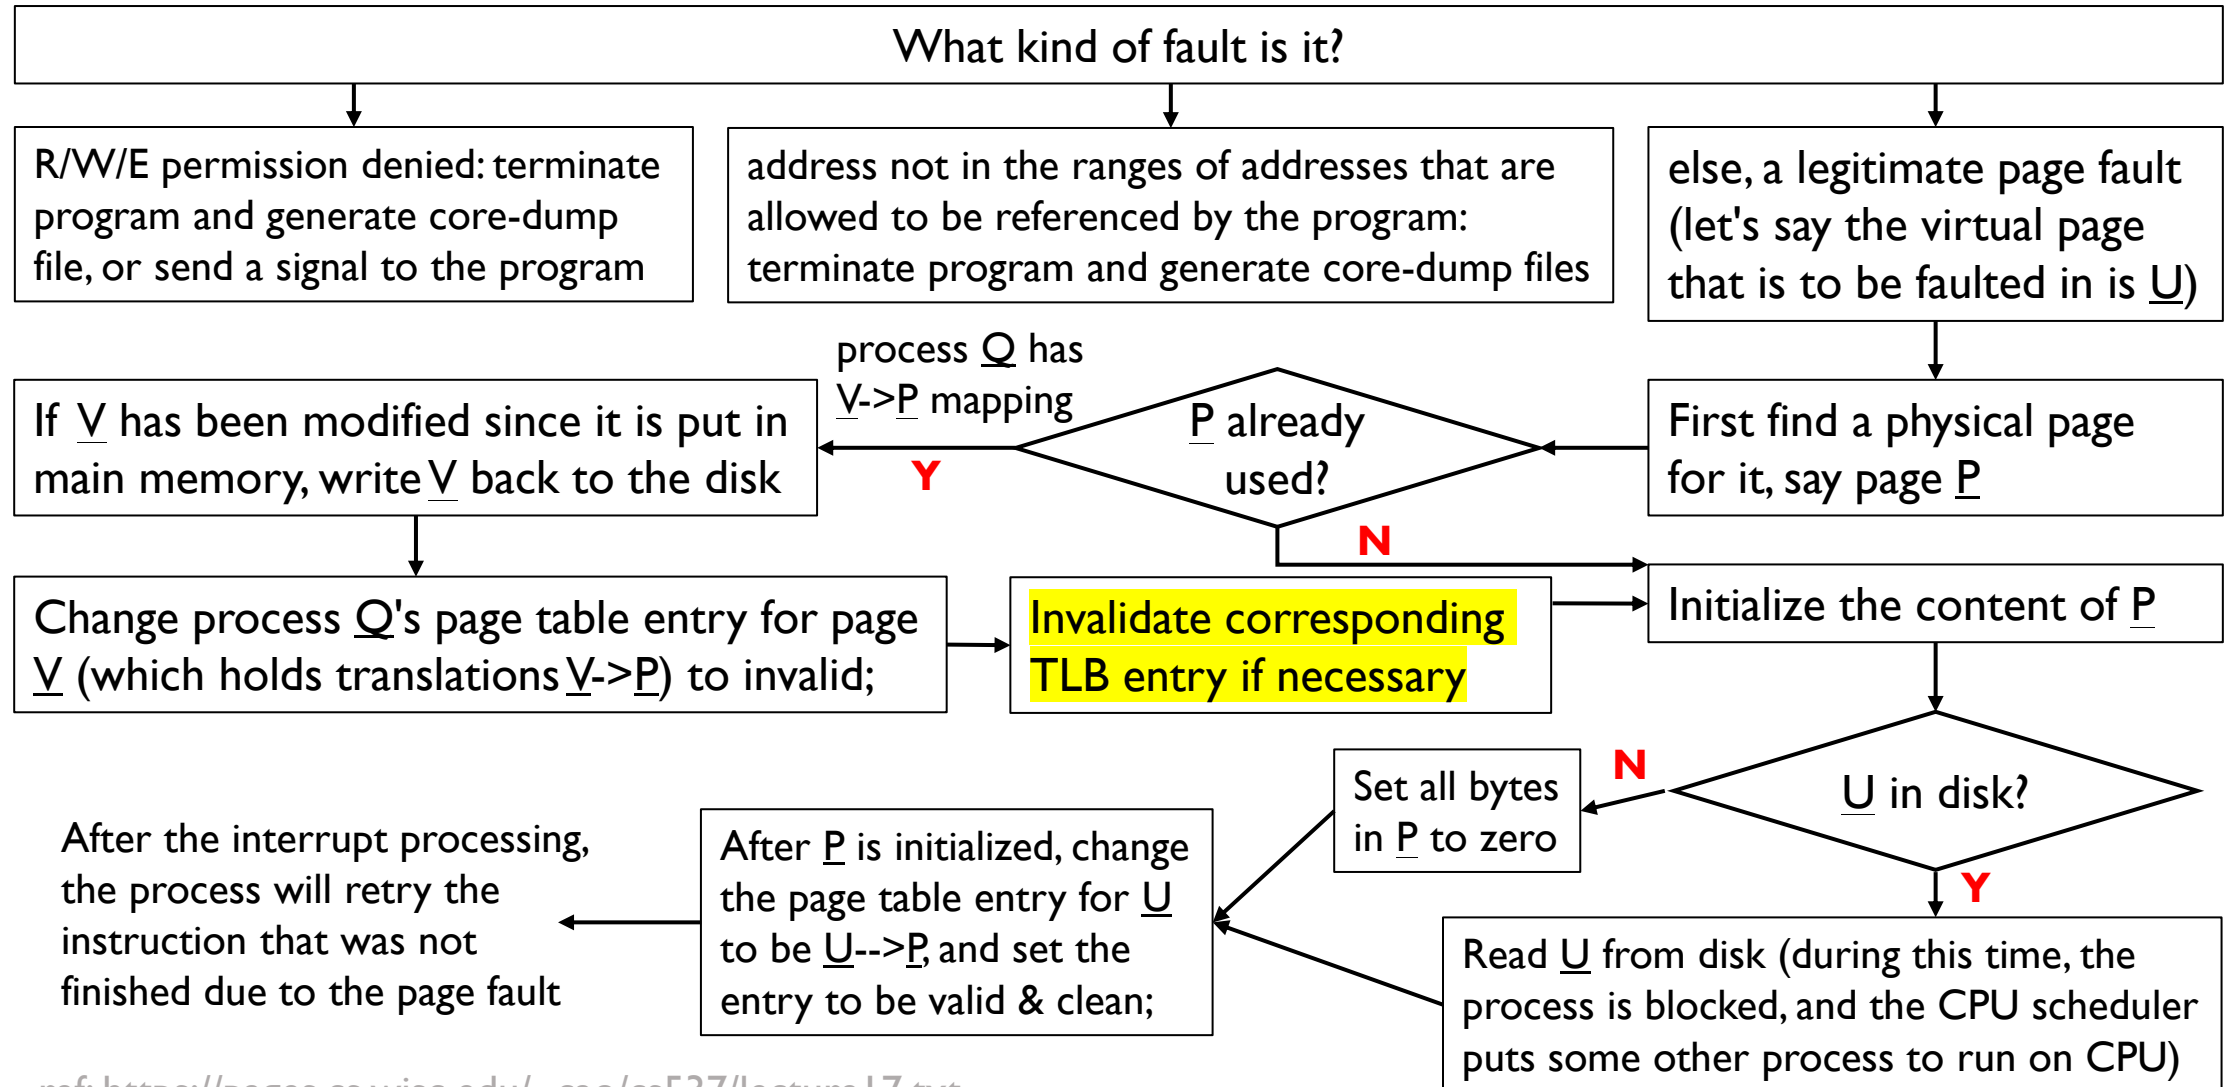
Detailed Page Fault Process

Before Page Fault (done by hardware)



Detailed Page Fault Process

Handling Page Fault (done by hardware)



The Dirty Bit

- How does OS know which pages have been modified?
 - Assuming every page has been modified is correct but inefficient
- The hardware tracks it with a dirty bit in page table entry
 - Initialized to 0
 - Set to 1 whenever there is a store instruction for the page
- The TLB also has a dirty bit
- Unix has a background thread to clean pages when it's too full

Allocating New Page Frame

- If there is an empty page, use it
- If there is no empty page
 - Select a page to evict
 - ☐ Need a lightweight policy
 - Find page table entries that point to the evicted page
 - ☐ Core map – an array that maps physical page frames back to the table entries
 - Set page table entry to invalid
 - ☐ TLB shutdown is needed. Why?
 - Copy back any changes to the evicted page
 - ☐ Write back
 - ☐ The same for application exit
 - ☐ **Dirty bit**

Page Eviction Policy

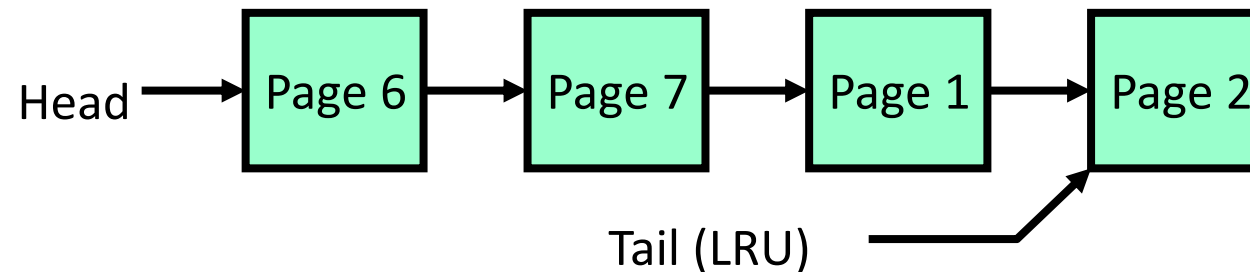
- Why do we care about Replacement Policy?
 - The cost of being wrong is high: must go to disk
 - Must keep important pages in memory, not toss them out
- **FIFO (First In, First Out)**
 - Throw out oldest page. Let every page live in memory for same amount of time.
 - Bad – throws out heavily used pages instead of infrequently used
- **MIN (Minimum):**
 - Replace page that won't be used for the longest time
 - Great, but can't really know future...
 - Makes good comparison case, however
- **RANDOM:**
 - Pick random page for every replacement
 - Typical solution for TLB's. Simple hardware
 - Pretty unpredictable – makes it hard to make real-time guarantees

Page Eviction Policy

- **LRU (Least Recently Used):**

- Replace page that hasn't been used for the longest time
- Programs have locality, so if something not used for a while, unlikely to be used in the near future.

- How to implement LRU? Use a list!



- On each use, remove page from list and place at head
- LRU page is at tail
- Problems with this scheme for paging?
 - Need to know immediately when each page is used, so we can change its position in list
 - Many instructions for each hardware access
- In practice, people **approximate** LRU

Page Eviction Policy


- Why we can implement LRU for TLB entry replacement, but not demand paging replacement?

Page Eviction Policy

- Why we can implement LRU for TLB entry replacement, but not demand paging replacement?
 - TLB is purely handled in hardware (MMU)
 - TLB has fewer entries (typically 16-512)

Page Eviction Policy


- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it



Page reference stream:

Page Eviction Policy


- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it



Page reference stream: 1

Page Eviction Policy


- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it



Page reference stream: 1

Page Eviction Policy


- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it



Page reference stream: 1 3

Page Eviction Policy


- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it



Page reference stream: 1 3

Page Eviction Policy


- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it



Page reference stream: 1 3 1

Page Eviction Policy


- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it



Page reference stream: 1 3 1 20

Page Eviction Policy


- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it



Page reference stream: 1 3 1 20

Page Eviction Policy


- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it



Page reference stream: 1 3 1 20

Page Eviction Policy


- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it



Page reference stream: 1 3 1 20 10

Page Eviction Policy


- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it



Page reference stream: 1 3 1 20 10 25

Page Eviction Policy


- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it



Page reference stream: 1 3 1 20 10 25

Page Eviction Policy


- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it



Page reference stream: 1 3 1 20 10 25

Page Eviction Policy


- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it



Page reference stream: 1 3 1 20 10 25 2

Page Eviction Policy


- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it



Page reference stream: 1 3 1 20 10 25 2

Page Eviction Policy

- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it




Page reference stream: 1 3 1 20 10 25 2

Page Eviction Policy

- **Clocking algorithm**: approximating LRU
- Implementation with the use bit
 - Initialized to 0 in page table
 - Set to 1 whenever there is a page access
- When we need to evict a page, we look at the page under the hand:
 - If its use bit = 1, clear it and move the hand, repeat;
 - If its use bit = 0, evict it

What if hand moving slowly? Good sign or bad sign?
What if hand moving quickly? Good sign or bad sign?



Page reference stream: 1 3 1 20 10 25 2

Nth Chance Version of Clock Algorithm

- **Nth chance algorithm:** Give page N chances
 - OS keeps counter per page: # sweeps
 - On page fault, OS checks use bit:
 - ❑ 1 → clear use and also clear counter (used in last sweep)
 - ❑ 0 → increment counter; if count=N, replace page
 - Means that clock hand has to sweep by N times without page being used before page is replaced
- How do we pick N?
 - Why pick large N? Better approximation to LRU
 - ❑ If N ~ 1K, really good approximation
 - Why pick small N? More efficient
 - ❑ Otherwise might have to look a long way to find free page
- What about dirty pages?
 - Takes extra overhead to replace a dirty page, so give dirty pages an extra chance before replacing?
 - Common approach:
 - ❑ Clean pages, use N=1
 - ❑ Dirty pages, use N=2 (and write back to disk when N=1)

Details of Clock Algorithms

- Which bits of a PTE entry are useful to us?
 - **Use**: set when page is referenced; cleared by clock algorithm
 - **Modified**: set when page is modified, cleared when page written to disk
 - **Valid**: ok for program to reference this page
 - **Read-only**: ok for program to read page, but not modify
 - ☐ For example for catching modifications to code pages!
- Do we really need hardware-supported “modified” bit?
 - No. Can emulate it (BSD Unix) using read-only bit
 - ☐ Initially, mark all pages as read-only, even data pages
 - ☐ On write, trap to OS. OS sets software “modified” bit, and marks page as read-write.
 - ☐ Whenever page comes back in from disk, mark read-only

Allocation of Page Frames

- How do we allocate memory among different processes?
 - Does every process get the same fraction of memory? Different fractions?
 - Should we completely swap some processes out of memory?
- Each process needs *minimum* number of pages
 - Want to make sure that all processes **that are loaded into memory** can make forward progress
 - Example: IBM 370 – 6 pages to handle SS MOVE instruction:
 - ☐ instruction is 6 bytes, might span 2 pages
 - ☐ 2 pages to handle *from*
 - ☐ 2 pages to handle *to*
- Possible Replacement Scopes:
 - **Global replacement** – process selects replacement frame from set of all frames; one process can take a frame from another
 - **Local replacement** – each process selects from only its own set of allocated frames

Allocation of Page Frames

- Self-paging (自分页): each process is responsible for managing its own page faults and memory allocation, rather than relying on a global operating system-wide policy.
- Global page management
 - each process/user is assigned its fair share of page frames using max-min scheduling algorithm
 - when memory is full, the page eviction happens to the process with the most allocated memory.
 - ❑ Avoid malicious attackers that wants as much as resources

Summary

- To support demand paging, what do CPU/OS contribute?
 - CPU: memory management (MMU), a few bits in page table entry, etc
 - OS: page table manipulation, eviction strategy, page fault handler, etc

Advanced: Android Memory Management

- How does Android (or other mobile OSes) handle memory inefficiency, e.g., too many apps opened?
 - Strategy #1: swapping (demand paging)
 - (primary) Strategy #2: low memory killer (LMK)
 - ❑ vs. out-of-memory (OOM) killer in Linux
 - Android prefers the 2nd one, because:
 - ❑ Flash memory has limited write endurance.
 - ❑ Disk I/O is generally slower and consumes more power compared to RAM access.
 - ❑ Responsive time is more important on mobile apps