

# Operating Systems

## Lecture 12

### Readers/Writers and Deadlock

Prof. Mengwei Xu



# Goals for Today

---

- Readers/Writers Lock
- Deadlock




# Goals for Today

---

- Readers/Writers Lock
- Deadlock

# Readers/Writers Problem



- Motivation: Consider a shared database
  - Two classes of users:
    - Readers – never modify database
    - Writers – read and modify database
  - Is using a single lock on the whole database sufficient?
    - Like to have many readers at the same time
    - Only one writer at a time

# Basic Readers/Writers Solution

---

- Correctness Constraints:
  - Readers can access database when no writers
  - Writers can access database when no readers or writers
  - Only one thread manipulates state variables at a time
- Basic structure of a solution:
  - **Reader()**  
Wait until no writers  
Access data base  
Check out - wake up a waiting writer
  - **Writer()**  
Wait until no active readers or writers  
Access database  
Check out - wake up waiting readers or writer
  - State variables (Protected by a lock called “lock”):
    - ❑ int AR: Number of active readers; initially = 0
    - ❑ int WR: Number of waiting readers; initially = 0
    - ❑ int AW: Number of active writers; initially = 0
    - ❑ int WW: Number of waiting writers; initially = 0
    - ❑ Condition okToRead = NIL
    - ❑ Condition okToWrite = NIL

# Code for a Reader

```
Reader() {  
    // First check self into system  
    lock.Acquire();  
  
    while ((AW + WW) > 0) { // Is it safe to read?  
        WR++; // No. Writers exist  
        okToRead.wait(&lock); // Sleep on cond var  
        WR--; // No longer waiting  
    }  
    AR++; // Now we are active!  
    lock.release();  
  
    // Perform actual read-only access  
    AccessDatabase(ReadOnly);  
  
    // Now, check out of system  
    lock.Acquire();  
    AR--; // No longer active  
    if (AR == 0 && WW > 0) // No other active readers  
        okToWrite.signal(); // Wake up one writer  
    lock.Release();  
}
```

Why release lock here?

# Code for a Writer

```

Writer() {
    // First check self into system
    lock.Acquire();

    while ((AW + AR) > 0) { // Is it safe to write?
        WW++; // No. Active users exist
        okToWrite.wait(&lock); // Sleep on cond var
        WW--; // No longer waiting
    }

    AW++; // Now we are active!
    lock.release();

    // Perform actual
    AccessDatabase(Request);

    // Now, check out
    lock.Acquire();
    AW--;
    if (WW > 0) {
        okToWrite.signal(); // Give priority to writers
    } else if (WR > 0) {
        okToRead.broadcast(); // Otherwise, wake reader
    }
    lock.Release();
}

```

Why broadcast() here instead of signal()?

Why Give priority to writers?

# Simulation of Readers/Writers Solution

---

- Use an example to simulate the solution
- Consider the following sequence of operators:
  - R1, R2, W1, R3
- Initially:  $AR = 0$ ,  $WR = 0$ ,  $AW = 0$ ,  $WW = 0$

# Simulation of Readers/Writers Solution

---

- R1 comes along
- AR = 0, WR = 0, AW = 0, WW = 0

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R1 comes along
- AR = 0, WR = 0, AW = 0, WW = 0

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R1 comes along
- AR = 1, WR = 0, AW = 0, WW = 0

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();
}

AccessDbase(ReadOnly);

lock.Acquire();
AR--;
if (AR == 0 && WW > 0)
    okToWrite.signal();
lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R1 comes along
- AR = 1, WR = 0, AW = 0, WW = 0

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();
```

```
AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R1 comes along
- AR = 1, WR = 0, AW = 0, WW = 0

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();
}
```

**AccessDbase(ReadOnly)**

```
lock.Acquire();
AR--;
if (AR == 0 && WW > 0)
    okToWrite.signal();
lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R2 comes along
- AR = 1, WR = 0, AW = 0, WW = 0

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R2 comes along
- AR = 1, WR = 0, AW = 0, WW = 0

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R2 comes along
- AR = 2, WR = 0, AW = 0, WW = 0

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();
}

AccessDbase(ReadOnly);

lock.Acquire();
AR--;
if (AR == 0 && WW > 0)
    okToWrite.signal();
lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R2 comes along
- AR = 2, WR = 0, AW = 0, WW = 0

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();
```

```
AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R2 comes along
- AR = 2, WR = 0, AW = 0, WW = 0

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();
}
```

**AccessDbase (ReadOnly)**

```
lock.Acquire();
AR--;
if (AR == 0 && WW > 0)
    okToWrite.signal();
}
```

1 Assume readers take a while to access database  
Situation: Locks released, only AR is non-zero

# Simulation of Readers/Writers Solution

---

- W1 comes along (R1 and R2 are still accessing dbase)
- AR = 2, WR = 0, AW = 0, WW = 0

```
Writer() {
    lock.Acquire();
    while ((AW + AR) > 0) { // Is it safe to write?
        WW++;
        okToWrite.wait(&lock); // No. Active users exist
        WW--;
    }
    AW++;
    lock.release();

    AccessDatabase(ReadWrite);

    lock.Acquire();
    AW--;
    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- W1 comes along (R1 and R2 are still accessing dbase)
- AR = 2, WR = 0, AW = 0, WW = 0

```
Writer() {
    lock.Acquire();
    while ((AW + AR) > 0) { // Is it safe to write?
        WW++;
        okToWrite.wait(&lock); // No. Active users exist
        WW--;
    }
    AW++;
    lock.release();

    AccessDatabase(ReadWrite);

    lock.Acquire();
    AW--;
    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- W1 comes along (R1 and R2 are still accessing dbase)
- AR = 2, WR = 0, AW = 0, WW = 1

```
Writer() {
    lock.Acquire();
    while ((AW + AR) > 0) { // Is it safe to write?
        WW++;
        okToWrite.wait(&lock); // No. Active users exist
        WW--;
    }
    AW++;
    lock.release();

    AccessDatabase(ReadWrite);

    lock.Acquire();
    AW--;
    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

- W1 comes along (R1 and R2 are still accessing dbase)
- AR = 2, WR = 0, AW = 0, WW = 1

```
Writer() {
    lock.Acquire();
    while ((AW + AR) > 0) { // Is it safe to write?
        WW++;
        okToWrite.wait(&lock); // No. Active users exist
        WW--;
    }
    AW++;
    lock.release();

    AccessDatabase(ReadWrite);

    lock.Acquire();
    AW--;
    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    lock.Release();
}
```

W1 cannot start because of readers, so goes to sleep

# Simulation of Readers/Writers Solution

---

- R3 comes along (R1, R2 accessing dbase, W1 waiting)
- AR = 2, WR = 0, AW = 0, WW = 1

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R3 comes along (R1, R2 accessing dbase, W1 waiting)
- AR = 2, WR = 0, AW = 0, WW = 1

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R3 comes along (R1, R2 accessing dbase, W1 waiting)
- AR = 2, WR = 1, AW = 0, WW = 1

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++; // No. Writers exist
        okToRead.wait(&lock); // Sleep on cond var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

- R3 comes along (R1, R2 accessing dbase, W1 waiting)
- AR = 2, WR = 1, AW = 0, WW = 1

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // No. Writers exist
        WR--;
        // Sleep on cond var
        // No longer waiting
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
}
```

Status:

- R1 and R2 still reading
- W1 and R3 waiting on okToWrite and okToRead, respectively

# Simulation of Readers/Writers Solution

---

- R2 finishes (R1 accessing dbase, W1, R3 waiting)
- AR = 2, WR = 1, AW = 0, WW = 1

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire(); // Line highlighted with a blue box
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R2 finishes (R1 accessing dbase, W1, R3 waiting)
- AR = 1, WR = 1, AW = 0, WW = 1

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R2 finishes (R1 accessing dbase, W1, R3 waiting)
- AR = 1, WR = 1, AW = 0, WW = 1

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R2 finishes (R1 accessing dbase, W1, R3 waiting)
- AR = 1, WR = 1, AW = 0, WW = 1

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R1 finishes (W1, R3 waiting)
- AR = 1, WR = 1, AW = 0, WW = 1

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire(); // Line highlighted with a blue box
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R1 finishes (W1, R3 waiting)
- AR = 0, WR = 1, AW = 0, WW = 1

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R1 finishes (W1, R3 waiting)
- AR = 0, WR = 1, AW = 0, WW = 1

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

- R1 finishes (W1, R3 waiting)
- AR = 0, WR = 1, AW = 0, WW = 1

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

All reader finished, signal writer – note, R3 still waiting

# Simulation of Readers/Writers Solution

- W1 gets signal (R3 still waiting)
- AR = 0, WR = 1, AW = 0, WW = 1

```
Writer() {
    lock.Acquire();
    while ((AW + AR) > 0) { // Is it safe to write?
        WW++;
        okToWrite.wait(&lock); // No. Active users exist
        WW--;
        // Sleep on cond var
        // No longer waiting
    }
    Got signal
    from R1
    lock.release();
    AccessDbase(ReadWrite);
}

lock.Acquire();
AW--;
if (WW > 0) {
    okToWrite.signal();
} else if (WR > 0) {
    okToRead.broadcast();
}
lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- W1 gets signal (R3 still waiting)
- AR = 0, WR = 1, AW = 0, WW = 0

```
Writer() {
    lock.Acquire();
    while ((AW + AR) > 0) { // Is it safe to write?
        WW++;
        okToWrite.wait(&lock); // No. Active users exist
        WW--; // Sleep on cond var
    } // No longer waiting
    AW++;
    lock.release();

    AccessDbase(ReadWrite);

    lock.Acquire();
    AW--;
    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- W1 gets signal (R3 still waiting)
- AR = 0, WR = 1, AW = 1, WW = 0

```
Writer() {
    lock.Acquire();
    while ((AW + AR) > 0) { // Is it safe to write?
        WW++;
        okToWrite.wait(&lock); // No. Active users exist
        WW--;
    }
    AW++;
    lock.release();
}

AccessDbase(ReadWrite);

lock.Acquire();
AW--;
if (WW > 0) {
    okToWrite.signal();
} else if (WR > 0) {
    okToRead.broadcast();
}
lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- W1 gets signal (R3 still waiting)
- AR = 0, WR = 1, AW = 1, WW = 0

```
Writer() {
    lock.Acquire();
    while ((AW + AR) > 0) { // Is it safe to write?
        WW++;
        okToWrite.wait(&lock); // No. Active users exist
        WW--;
    }
    AW++;
    lock.release();
}
```

**AccessDbase (ReadWrite) ;**

```
lock.Acquire();
AW--;
if (WW > 0) {
    okToWrite.signal();
} else if (WR > 0) {
    okToRead.broadcast();
}
lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- W1 gets signal (R3 still waiting)
- AR = 0, WR = 1, AW = 0, WW = 0

```
Writer() {
    lock.Acquire();
    while ((AW + AR) > 0) { // Is it safe to write?
        WW++;
        okToWrite.wait(&lock); // No. Active users exist
        WW--;
    }
    AW++;
    lock.release();

    AccessDbase(ReadWrite);

    lock.Acquire();
    AW--;
    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- W1 gets signal (R3 still waiting)
- AR = 0, WR = 1, AW = 0, WW = 0

```
Writer() {
    lock.Acquire();
    while ((AW + AR) > 0) { // Is it safe to write?
        WW++;
        okToWrite.wait(&lock); // No. Active users exist
        WW--;
    }
    AW++;
    lock.release();

    AccessDbase(ReadWrite);

    lock.Acquire();
    AW--;
    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- W1 gets signal (R3 still waiting)
- AR = 0, WR = 1, AW = 0, WW = 0

```
Writer() {
    lock.Acquire();
    while ((AW + AR) > 0) { // Is it safe to write?
        WW++;
        okToWrite.wait(&lock); // No. Active users exist
        WW--;
    }
    AW++;
    lock.release();

    AccessDbase(ReadWrite);

    lock.Acquire();
    AW--;
    if (WW > 0) {
        okToWrite.signal();
    } else if (WR > 0) {
        okToRead.broadcast();
    }
    lock.Release();
}
```

No waiting writer, signal reader R3

# Simulation of Readers/Writers Solution

- R1 finishes (W1, R3 waiting)
- AR = 0, WR = 1, AW = 0, WW = 0

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // No. Writers exist
        // Sleep on cond var
        WR--;
        // No longer waiting
    }
    Got signal
    from W1; // Now we are active!
    lock.release();
}

AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R1 finishes (W1, R3 waiting)
- AR = 0, WR = 0, AW = 0, WW = 0

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--; // No longer waiting
    }
    AR++; // Now we are active!
    lock.release();
}

AccessDbase(ReadOnly);

lock.Acquire();
AR--;
if (AR == 0 && WW > 0)
    okToWrite.signal();
lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R1 finishes (W1, R3 waiting)
- AR = 0, WR = 0, AW = 0, WW = 0

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();
}
```

**AccessDbase(ReadOnly)**

```
lock.Acquire();
AR--;
if (AR == 0 && WW > 0)
    okToWrite.signal();
lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R1 finishes (W1, R3 waiting)
- AR = 0, WR = 0, AW = 0, WW = 0

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire(); // Line highlighted with a blue box
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

# Simulation of Readers/Writers Solution

---

- R1 finishes (W1, R3 waiting)
- AR = 0, WR = 0, AW = 0, WW = 0

```
Reader() {
    lock.Acquire();
    while ((AW + WW) > 0) { // Is it safe to read?
        WR++;
        okToRead.wait(&lock); // Sleep on cond var
        WR--;
    }
    AR++; // Now we are active!
    lock.release();

    AccessDbase(ReadOnly);

    lock.Acquire();
    AR--;
    if (AR == 0 && WW > 0)
        okToWrite.signal();
    lock.Release();
}
```

DONE!

# Read/Writer Questions

```
Reader() {  
    // check into system  
    lock.Acquire();  
    while ((AW + WW) > 0) {  
        WR++;  
        okToRead.wait(&lock);  
        WR--;  
    }  
    AR++;  
    lock.release();  
  
    // read-only  
    AccessDbase();  
  
    // check out  
    lock.Acquire();  
    AR--;  
    if (AR == 0 && WW > 0)  
        okToWrite.signal();  
    lock.Release();  
}
```

What if we  
remove this  
line?

```
Writer() {  
    // check into system  
    lock.Acquire();  
    while ((AW + AR) > 0) {  
        WW++;  
        okToWrite.wait(&lock);  
        WW--;  
    }  
    AW++;  
    lock.release();  
  
    // read/write access  
    AccessDbase(ReadWrite);  
  
    // check out of system  
    lock.Acquire();  
    AW--;  
    if (WW > 0){  
        okToWrite.signal();  
    } else if (WR > 0){  
        okToRead.broadcast();  
    }  
    lock.Release();  
}
```

# Read/Writer Questions

```
Reader() {  
    // check into system  
    lock.Acquire();  
    while ((AW + WW) > 0) {  
        WR++;  
        okToRead.wait(&lock);  
        WR--;  
    }  
    AR++;  
    lock.release();  
  
    // read-only  
    AccessDbase();  
  
    // check out  
    lock.Acquire();  
    AR--;  
    if (AR == 0 && WR > 0)  
        okToWrite.broadcast();  
    lock.Release();  
}
```

What if we turn  
signal to  
broadcast?

```
Writer() {  
    // check into system  
    lock.Acquire();  
    while ((AW + AR) > 0) {  
        WW++;  
        okToWrite.wait(&lock);  
        WW--;  
    }  
    AW++;  
    lock.release();  
  
    // read/write access  
    AccessDbase(ReadWrite);  
  
    // check out of system  
    lock.Acquire();  
    AW--;  
    if (WW > 0){  
        okToWrite.signal();  
    } else if (WR > 0){  
        okToRead.broadcast();  
    }  
    lock.Release();  
}
```

# Read/Writer Questions

---

```
Reader() {  
    // check into system  
    lock.Acquire();  
    while ((AW + WW) > 0) {  
        WR++;  
        okContinue.wait(&lock);  
        WR--;  
    }  
    AR++;  
    lock.release();  
  
    // read-only access  
    AccessDbase(ReadOnly);  
  
    // check out of system  
    lock.Acquire();  
    AR--;  
    if (AR == 0 && WW > 0)  
        okContinue.signal();  
    lock.Release();  
}  
  
Writer() {  
    // check into system  
    lock.Acquire();  
    while ((AW + AR) > 0) {  
        WW++;  
        okContinue.wait(&lock);  
        WW--;  
    }  
    AW++;  
    lock.release();  
  
    // read/write access  
    AccessDbase(ReadWrite);  
  
    // check out of system  
    lock.Acquire();  
    AW--;  
    if (WW > 0){  
        okContinue.signal();  
    } else if (WR > 0) {  
        okContinue.broadcast();  
    }  
    lock.Release();  
}
```

What if we turn **okToWrite** and **okToRead** into **okContinue**?

# Read/Writer Questions

```
Reader() {  
    // check into system  
    lock.Acquire();  
    while ((AW + WW) > 0) {  
        WR++;  
        okContinue.wait(&lock);  
        WR--;  
    }  
    AR++;  
    lock.release();  
  
    // read-only access  
    AccessDbase(ReadOnly);  
  
    // check out of system  
    lock.Acquire();  
    AR--;  
    if (AR == 0 && WW > 0)  
        okContinue.signal();  
    lock.Release();  
}  
  
Writer() {  
    // check into system  
    lock.Acquire();  
    while ((AW + AR) > 0) {  
        WW++;  
        okContinue.wait(&lock);  
        WW--;  
    }  
    AW++;  
    lock.release();  
  
    // read/write access  
    AccessDbase(ReadWrite);  
  
    // check out of system  
    lock.Acquire();  
    AW--;  
    if (WW > 0){  
        okContinue.signal();  
    } else if (WR > 0){  
        okContinue.broadcast();  
    }  
    lock.Release();  
}
```

- R1 arrives
- W1, R2 arrive while R1 still reading → W1 and R2 wait for R1 to finish
- Assume R1's signal is delivered to R2 (not W1)

# Read/Writer Questions

```
Reader() {  
    // check into system  
    lock.Acquire();  
    while ((AW + WW) > 0) {  
        WR++;  
        okContinue.wait(&lock);  
        WR--;  
    }  
    AR++;  
    lock.release();  
  
    // read-only access  
    AccessDbase(ReadOnly);  
  
    // check out of system  
    lock.Acquire();  
    AR--;  
    if (AR == 0 && WW > 0)  
        okContinue.broadcast();  
    lock.Release();  
}  
  
Writer() {  
    // check into system  
    lock.Acquire();  
    while ((AW + AR) > 0) {  
        WW++;  
        okContinue.wait(&lock);  
        WW--;  
    }  
    AW++;  
    lock.release();  
  
    // read/write access  
    AccessDbase(ReadWrite);  
  
    // check out of system  
    lock.Acquire();  
    AW--;  
    if (WW > 0){  
        okContinue.signal();  
    } else if (WR > 0) {  
        okContinue.broadcast();  
    }  
    lock.Release();  
}
```

Need to change to broadcast!

# Implementing RLock

---

- Let's wrap the code into a RLock class

```
RLock* rwlock;  
  
rwlock->startRead();  
// Read shared data  
rwlock->doneRead();  
  
rwlock->startWrite();  
// Write shared data  
rwlock->startRead();
```

# Implementing RWLock

```
class RWLock {  
    Lock lock;  
    CV canRead;  
    CV canWrite;  
    int AR, AW, WR, WW;  
}
```

```
void RWLock::startRead() {  
    lock.acquire();  
    WR++;  
    while ((AW + WW > 0)) {  
        canRead.Wait(&lock);  
    }  
    WR--;  
    AR++;  
    lock.release();  
}
```

```
void RWLock::doneRead() {  
    lock.acquire();  
    AR--;  
    if ((AR == 0) && (WW > 0)) {  
        canWrite.signal();  
    }  
    lock.release();  
}
```

# Implementing RWLock

```
class RWLock {  
    Lock lock;  
    CV canRead;  
    CV canWrite;  
    int AR, AW, WR, WW;  
}
```

```
void RWLock::startWrite() {  
    lock.acquire();  
    WW++;  
    while ((AW + AR > 0)) {  
        canWrite.Wait(&lock);  
    }  
    WW--;  
    AW++;  
    lock.release();  
}
```

```
void RWLock::doneWrite() {  
    lock.acquire();  
    AW--;  
    assert(AW == 0);  
    if (WW > 0) {  
        canWrite.signal();  
    }  
    else {  
        canRead.broadcast();  
    }  
    lock.release();  
}
```



# Goals for Today

---

- Readers/Writers Lock
- Deadlock

# Deadlock

---

- Deadlock (死锁): a cycle of waiting among a set of threads, where each thread waits for some other thread in the cycle to take some action.
- A simple case: mutually recursive locking

// Thread A

```
lock1.acquire();  
lock2.acquire();  
lock2.release();  
lock1.release();
```

// Thread B

```
lock2.acquire();  
lock1.acquire();  
lock1.release();  
lock2.release();
```

# Deadlock

---

- Deadlock (死锁): a cycle of waiting among a set of threads, where each thread waits for some other thread in the cycle to take some action.
- Another example with 2 locks and 1 condition variable

// Thread A

```
lock1.acquire();
lock2.acquire();
while (need to wait) {
    cv.wait(&lock2);
}
lock2.release();
lock1.release();
```

// Thread B

```
lock1.acquire();
lock2.acquire();
cv.signal();
lock2.release();
lock1.release();
```

# Deadlock

- Deadlock (死锁): a cycle of waiting among a set of threads, where each thread waits for some other thread in the cycle to take some action.
- Another example with 2 locks and 1 condition variable

// Thread A

```
lock1.acquire();
lock2.acquire();
while (need to wait) {
    cv.wait(&lock2);
}
lock2.release();
lock1.release();
```


// Thread B

```
lock1.acquire();
lock2.acquire();
cv.signal();
lock2.release();
lock1.release();
```

Any deadlock?

# Starvation vs Deadlock

- Starvation vs. Deadlock
  - Starvation: thread waits indefinitely
    - Example, low-priority thread waiting for resources constantly in use by high-priority threads
  - Deadlock: circular waiting for resources
    - Thread A owns Res 1 and is waiting for Res 2
    - Thread B owns Res 2 and is waiting for Res 1
  - Deadlock  $\Rightarrow$  Starvation but not vice versa
    - Starvation can end (but doesn't have to)
    - Deadlock can't end without external intervention




# Bridge Crossing Example

- Each segment of road can be viewed as a resource
  - Car must own the segment under them
  - Must acquire segment that they are moving into
- For bridge: must acquire both halves
  - Traffic only in one direction at a time
  - Problem occurs when two cars in opposite directions on bridge: each acquires one segment and needs next
- If a deadlock occurs, it can be resolved if one car backs up (preempt resources and rollback)
  - Several cars may have to be backed up
- Starvation is possible
  - East-going traffic really fast  $\Rightarrow$  no one goes west



# Dining Philosophers Problem

- Dining Philosophers Problem (哲学家进餐问题)
  - For example: 5 philosophers, 5 plate, and 5 chopsticks
  - When a philosopher thinking, he holds nothing
  - When a philosopher wants to eat, he first picks up the left chopstick, and then the right chopstick. After eating, he puts down both chopsticks.
- Stuck when everyone holds the left chopstick
- A general case of mutually recursive locking



# Conditions for Deadlock

---

- Deadlock not always deterministic – Example 2 mutexes:

Thread A

```
x.P();  
y.P();  
y.V();  
x.V();
```

Thread B

```
y.P();  
x.P();  
x.V();  
y.V();
```

- Deadlock won't always happen with this code
  - Have to have exactly the right timing ("wrong" timing?)
  - So you release a piece of software, and you tested it, and there it is, controlling a nuclear power plant...
- Deadlocks occur with multiple resources
  - Means you can't decompose the problem
  - Can't solve deadlock for each resource independently
- Example: System with 2 disk drives and two threads
  - Each thread needs 2 disk drives to function
  - Each thread gets one disk and waits for another one

# Four requirements for Deadlock

---

- Mutual exclusion
  - Only one thread at a time can use a resource.
- Hold and wait
  - Thread holding at least one resource is waiting to acquire additional resources held by other threads
- No preemption
  - Resources are released only voluntarily by the thread holding the resource, after thread is finished with it
- Circular wait
  - There exists a set  $\{T_1, \dots, T_n\}$  of waiting threads
    - $T_i$  is waiting for a resource that is held by  $T_{i+1}$

# Four requirements for Deadlock

---

- Mutual exclusion
  - Only one thread at a time can use a resource.
  - Each chopstick can be held by a single philosopher at a time
- Hold and wait
  - Thread holding at least one resource is waiting to acquire additional resources held by other threads
  - When a philosopher needs to wait for a chopstick, he continues to hold onto any chopsticks he has already picked up
- No preemption
  - Resources are released only voluntarily by the thread holding the resource, after thread is finished with it
  - Once a philosopher picks up a chopstick, he does not release it until he is done eating.
- Circular wait
  - There exists a set  $\{T_1, \dots, T_n\}$  of waiting threads
    - $T_i$  is waiting for a resource that is held by  $T_{i+1}$
  - Everyone is holding the left chopstick but waiting for the right one.

# Methods for Handling Deadlocks

---

- Allow system to enter deadlock and then recover
  - Requires deadlock detection algorithm
  - Some technique for forcibly preempting resources and/or terminating tasks
- Ensure that system will *never* enter a deadlock
  - Need to monitor all lock acquisitions
  - Selectively deny those that *might* lead to deadlock
- Ignore the problem and pretend that deadlocks never occur in the system
  - Used by most operating systems, including UNIX

# Preventing Deadlocks

---

1. No circular wait
2. No hold-and-wait
3. No mutual exclusion
4. Smart scheduling
  - banking algorithm

# Preventing Deadlocks

---

1. No circular wait
2. No hold-and-wait
3. No mutual exclusion
4. Smart scheduling
  - banking algorithm

# Removing Circular Wait

- Just make sure all locks acquired in the same order!
  - Total ordering
  - Partial ordering
  - An excellent example: memory mapping code in Linux

<https://github.com/torvalds/linux/blob/master/mm/filemap.c>

```
/*
 * Lock ordering:
 *
 * ->i_mmap_rwsem          (truncate_pagecache)
 *   ->private_lock          (_free_pte->block_dirty_folio)
 *   ->swap_lock              (exclusive_swap_page, others)
 *   ->i_pages lock
 *
 * ->i_rwsem                (acquired by fs in truncate path)
 *   ->invalidate_lock
 *   ->i_mmap_rwsem          (truncate->unmap_mapping_range)
 *
 * ->mmap_lock
 *   ->i_mmap_rwsem
 *   ->page_table_lock or pte_lock (various, mainly in memory.c)
 *   ->i_pages lock          (arch-dependent flush_dcache_mmap_lock)
 *
 * ->mmap_lock
 *   ->invalidate_lock
 *   ->lock_page              (filemap_fault)
 *   ->lock_page              (filemap_fault, access_process_vm)
 *
 * ->i_rwsem                (generic_perform_write)
 *   ->mmap_lock              (fault_in_readable->do_page_fault)
 *
 * bdi->wb.list_lock
 *   ->sb_lock                (fs/fs-writeback.c)
 *   ->i_pages lock            (_sync_single_inode)
 *
 * ->i_mmap_rwsem
 *   ->anon_vma.lock          (vma_merge)
 *
 * ->anon_vma.lock
 *   ->page_table_lock or pte_lock (anon_vma_prepare and various)
 *
 * ->page_table_lock or pte_lock
 *   ->swap_lock                (try_to_unmap_one)
 *   ->private_lock              (try_to_unmap_one)
 *   ->i_pages lock              (try_to_unmap_one)
 *   ->lruvec->lru_lock        (follow_page_mask->mark_page_accessed)
 *   ->lruvec->lru_lock        (check_pte_range->folio_isolate_lru)
 *   ->private_lock              (folio_remove_rmap_pte->set_page_dirty)
 *   ->i_pages lock              (folio_remove_rmap_pte->set_page_dirty)
 *   ->bdi.wb->list_lock        (folio_remove_rmap_pte->set_page_dirty)
 *   ->inode->i_lock            (folio_remove_rmap_pte->set_page_dirty)
 *   ->memcg->move_lock         (folio_remove_rmap_pte->set_page_dirty)
 *   ->bdi.wb->list_lock        (folio_remove_rmap_pte->folio_memcg_lock)
 *   ->inode->i_lock            (zap_pte_range->set_page_dirty)
 *   ->private_lock              (zap_pte_range->set_page_dirty)
 *   ->private_lock              (zap_pte_range->block_dirty_folio)
 */

```

# Removing Circular Wait

---

- Just make sure all locks acquired in the same order!
  - Total ordering
  - Partial ordering

```
func(mutex_t *m1, mutex_t *m2)
```

How to guarantee the ordering in func? Think about this case:

```
In Thread A: func(L1, L2)  
In Thread B: func(L2, L1)
```

# Removing Circular Wait

- Just make sure all locks acquired in the same order!
  - Total ordering
  - Partial ordering
- Enforce lock ordering by lock address

```
if (m1 > m2) { // grab in high-to-low address order
    pthread_mutex_lock(m1);
    pthread_mutex_lock(m2);
} else {
    pthread_mutex_lock(m2);
    pthread_mutex_lock(m1);
}
// Code assumes that m1 != m2 (not the same lock)
```

```
func(mutex_t *m1, mutex_t *m2)
```

How to guarantee the ordering in func? Think about this case:

In Thread A: func(L1, L2)

In Thread B: func(L2, L1)

# Preventing Deadlocks

---

1. No circular wait
  - Cons: needs careful design and programming from developers.
2. No hold-and-wait
3. No mutual exclusion
4. Smart scheduling
  - banking algorithm

# Preventing Deadlocks

---

1. No circular wait
2. **No hold-and-wait**
3. No mutual exclusion
4. Smart scheduling
  - banking algorithm

# Preventing Hold and Wait

---

- Just use another lock to lock the locks

```
pthread_mutex_lock(prevention); // begin acquisition
pthread_mutex_lock(L1);
pthread_mutex_lock(L2);
...
pthread_mutex_unlock(prevention); // end
```

# Preventing Deadlocks

---

1. No circular wait
2. No hold-and-wait
  - Cons: must know which locks will be used beforehand; concurrency decreased.
3. No mutual exclusion
4. Smart scheduling
  - banking algorithm

# Preventing Deadlocks

---

1. No circular wait
2. No hold-and-wait
3. **No mutual exclusion**
4. Smart scheduling
  - banking algorithm

# Preventing Mutual Exclusion

---

- Design lock-free (or wait-free) data structures and algorithms using powerful hardware instructions

```
int CompareAndSwap(int *address, int expected,  
int new) {  
    if (*address == expected) {  
        *address = new;  
        return 1; // success  
    }  
    return 0; // failure  
}
```

# Preventing Mutual Exclusion

---

- Using **CompareAndSwap** to implement “increment a value by n”.

```
void AtomicIncrement(int *value, int n) {  
    do {  
        int old = *value;  
    } while (CompareAndSwap(value, old, old + n)==0);  
}
```

# Preventing Mutual Exclusion

---

- Using **CompareAndSwap** to implement “insert an element to a list head”.

```
// without deadlock prevention
void insert(int value) {
    node_t *n = malloc(sizeof(node_t));
    assert(n != NULL);
    n->value = value;
    n->next = head;
    head = n;
}
```

# Preventing Mutual Exclusion

---

- Using **CompareAndSwap** to implement “insert an element to a list head”.

```
// with deadlock prevention
void insert(int value) {
    node_t *n = malloc(sizeof(node_t));
    assert(n != NULL);
    n->value = value;
    do {
        n->next = head;
    } while (CompareAndSwap(&head, n->next, n) == 0);
}
```

# Preventing Deadlocks

---

1. No circular wait
2. No hold-and-wait
3. No mutual exclusion
  - Cons: too complicated; hardware support needed (possibly performance degradation).
4. Smart scheduling
  - banking algorithm


# Preventing Deadlocks

---




1. No circular wait
2. No hold-and-wait
3. No mutual exclusion
4. **Smart scheduling**
  - banking algorithm

# Resource-Allocation Graph


- System Model
  - A set of Threads  $T_1, T_2, \dots, T_n$
  - Resource types  $R_1, R_2, \dots, R_m$   
*CPU cycles, memory space, I/O devices*
  - Each resource type  $R_i$  has  $W_i$  instances
  - Each thread utilizes a resource as follows:
    - ❑ Request () / Use () / Release ()
- Resource-Allocation Graph:
  - $V$  is partitioned into two types:
    - ❑  $T = \{T_1, T_2, \dots, T_n\}$ , the set threads in the system.
    - ❑  $R = \{R_1, R_2, \dots, R_m\}$ , the set of resource types in system
  - request edge – directed edge  $T_i \rightarrow R_j$
  - assignment edge – directed edge  $R_j \rightarrow T_i$




# Resource Allocation Graph Examples




# Resource Allocation Graph Examples



Simple Resource Allocation Graph



Allocation Graph With Deadlock



Allocation Graph With Cycle, but No Deadlock

# Deadlock Detection Algorithm

- Only one of each type of resource  $\Rightarrow$  look for loops
- More General Deadlock Detection Algorithm

- Let  $[X]$  represent an m-ary vector of non-negative integers (quantities of resources of each type):


$[\text{FreeResources}]$ : Current free resources each type  
 $[\text{Request}_X]$ : Current requests from thread  $X$   
 $[\text{Alloc}_X]$ : Current resources held by thread  $X$

- See if tasks can eventually terminate on their own

```

[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
    done = true
    Foreach node in UNFINISHED {
        if ([Requestnode] <= [Avail]) {
            remove node from UNFINISHED
            [Avail] = [Avail] + [Allocnode]
            done = false
        }
    }
} until(done)
  
```

- Nodes left in UNFINISHED  $\Rightarrow$  deadlocked




# What to do when detect deadlock?

---

- Terminate thread, force it to give up resources
  - In Bridge example, Godzilla picks up a car; hurls it into the river. Deadlock solved!
  - But, not always possible – killing a thread holding a mutex leaves world inconsistent
- Preempt resources without killing off thread
  - Take away resources from thread temporarily
  - Doesn't always fit with semantics of computation
- Roll back actions of deadlocked threads
  - Hit the rewind button, pretend last few minutes never happened
  - For bridge example, make one car roll backwards (may require others behind him)
  - Common technique in databases (transactions)
  - Of course, if you restart in exactly the same way, may reenter deadlock once again
- Many operating systems use other options

# Resource Requests over Time

- Applications usually don't know exactly when/what they're going to request
- Resources are taken/released over time



# Bankers Algorithm (银行家算法)

---

- What if you don't know the order/amount of requests ahead of time?
- Must assume some worst-case “max” resource needed by each process
- Toward right idea:
  - State maximum resource needs in advance
  - Allow particular thread to proceed if:  
$$(\text{available resources} - \# \text{requested}) \geq \text{max remaining that might be needed by any thread}$$
  - Invariant: At all times, every request would succeed
    - Really conservative! Let's do something better.


# Bankers Algorithm (银行家算法)

---

- Invariant: At all times, there exists some order of requests that would succeed.
- Key ideas
  - A thread states its maximum resource requirements, but acquires and releases resources incrementally as the thread executes.
  - The runtime system delays granting some requests to ensure that the system never deadlocks.


# Safe State and Unsafe State

- **Safe state:** for any possible sequence of resource requests, there is at least one safe sequence of processing the requests that eventually succeeds in granting all pending and future requests.
- **Unsafe state:** there is at least one sequence of future resource requests that leads to deadlock no matter what processing order is tried.
- **Deadlocked state:** the system has at least one deadlock.




# Safe State and Unsafe State

- **Safe state:** for any possible sequence of resource requests, there is at least one safe sequence of processing the requests that eventually succeeds in granting all pending and future requests.
  - A system *in a safe state controls its own destiny*: for any workload, it can avoid deadlock by delaying the processing of some requests.



# Safe State and Unsafe State

- **Unsafe state:** there is at least one sequence of future resource requests that leads to deadlock no matter what processing order is tried.
  - An unsafe state does not always lead to deadlock
  - However, as long as the system remains in an unsafe state, a bad workload or unlucky scheduling of requests can force it to deadlock.



# Bankers Algorithm (银行家算法)

---

- Invariant: At all times, there exists some order of requests that would succeed.
- The banker's algorithm delays any request that takes it from a safe to an unsafe state.

# Bankers Algorithm (银行家算法)

- Delay a request that takes us into unsafe state.
- How to implement this?
  - Allocate resources dynamically
    - Evaluate each request and grant if some ordering of threads is still deadlock free afterward
  - Use deadlock detection algorithm presented earlier:
    - BUT: Assume each process needs "max" resources to finish

```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
    done = true
    Foreach node in UNFINISHED {
        if ([Requestnode] <= [Avail]) {
            remove node from UNFINISHED
            [Avail] = [Avail] + [Allocnode]
            done = false
        }
    }
} until(done)
```

# Bankers Algorithm (银行家算法)

- Delay a request that takes us into unsafe state.
- How to implement this?
  - Allocate resources dynamically
    - Evaluate each request and grant if some ordering of threads is still deadlock free afterward
  - Use deadlock detection algorithm presented earlier:
    - BUT: Assume each process needs "max" resources to finish

```
[Avail] = [FreeResources]
Add all nodes to UNFINISHED
do {
    done = true
    Foreach node in UNFINISHED {
        if ([Maxnode] - [Allocnode] <= [Avail]) {
            remove node from UNFINISHED
            [Avail] = [Avail] + [Allocnode]
            done = false
        }
    }
} until(done)
```

Each process might need "max" resources in order to finish

# Bankers Algorithm (银行家算法)

---

- Delay a request that takes us into unsafe state.
- How to implement this?
  - Allocate resources dynamically
    - Evaluate each request and grant if some ordering of threads is still deadlock free afterward
  - Use deadlock detection algorithm presented earlier:
    - BUT: Assume each process needs "max" resources to finish
- Keeps system in a "SAFE" state, i.e. there exists a sequence  $\{T_1, T_2, \dots, T_n\}$  with  $T_1$  requesting all remaining resources, finishing, then  $T_2$  requesting all remaining resources, etc..
- vs. "Require all before starting", the Banker's algorithm allows the sum of maximum resource needs of all current threads to be greater than total resources

# Bankers Algorithm (银行家算法)

---

- **EXAMPLE: Page allocation with the Banker's Algorithm.**
  - Suppose we have a system with 8 pages of memory and three processes: A, B, and C, which need 4, 5, and 5 pages to complete, respectively.
- They take turns requesting one page each, and the system grants requests in order

# Bankers Algorithm (银行家算法)

- EXAMPLE: Page allocation with the Banker's Algorithm.
  - Suppose we have a system with 8 pages of memory and three processes: A, B, and C, which need 4, 5, and 5 pages to complete, respectively.
- They take turns requesting one page each, and the system grants requests in order

Process	Allocation											Oops! Deadlock!	
A	0	1	1	1	2	2	2	3	3	3	3	wait	wait
B	0	0	1	1	1	2	2	2	3	3	3	3	wait
C	0	0	0	1	1	1	2	2	2	2	2	wait	wait
Total	0	1	2	3	4	5	6	7	8	8	8	8	8

# Bankers Algorithm (银行家算法)


- EXAMPLE: Page allocation with the Banker's Algorithm.
  - Suppose we have a system with 8 pages of memory and three processes: A, B, and C, which need 4, 5, and 5 pages to complete, respectively.
- What if we use banker's algorithm?

Process	Allocation															
A	0	1	1	1	2	2	2	3	3	3	4	0	0	0	0	0
B	0	0	1	1	1	2	2	2	wait	wait	wait	wait	3	4	4	5
C	0	0	0	1	1	1	2	2	2	wait	wait	wait	3	3	wait	4
Total	0	1	2	3	4	5	6	7	7	7	8	4	6	7	7	8

Tasks successfully finished

# Banker's Algorithm Example

- Banker's algorithm with dining philosophers
  - “Safe” (won’t cause deadlock) if when try to grab chopstick either:
    - Not last chopstick
    - Is last chopstick but someone will have two afterwards
  - What if  $k$ -handed philosopher? Don’t allow if:
    - It’s the last one, no one would have  $k$
    - It’s 2<sup>nd</sup> to last, and no one would have  $k-1$
    - It’s 3<sup>rd</sup> to last, and no one would have  $k-2$
    - ...



# Preventing Deadlocks

---

1. No circular wait
2. No hold-and-wait
3. No mutual exclusion
4. Smart scheduling
  - banking algorithm
  - Cons: must know the entire set of tasks and their resource demands beforehand; concurrency decreased.
  - Only used in limited scenarios such as embedded system.

# Techniques for Preventing Deadlock

---

- Infinite resources
  - Include enough resources so that no one ever runs out of resources. Doesn't have to be infinite, just large
  - Give illusion of infinite resources (e.g. virtual memory)
- No Sharing of resources (totally independent threads)
  - Often true (most things don't depend on each other)
  - Not very realistic in general (can't guarantee)

# Deadlock Prevention – The Reality

---

- Deadlock Prevention is HARD
  - How many resources will each thread need?
  - How many total resources are there?
- Also Slow/Impractical
  - Matrix of resources/requirements could be big and dynamic
  - Re-evaluate on every request (even for small/non-contended)
  - Banker's algorithm assumes everyone asks for max
- REALITY
  - Most OSes don't bother
  - Programmers job to write deadlock-free programs (e.g. by ordering all resource requests).

# Homework

---

- Modify our RWLock implementation to use only one condition variable
- Implement Banker's Algorithm
  - Input-1: task number  $N$ , resource type number  $M$ ;
  - Input-2: resource amount: for each type:  $R_i$  where  $i=1-M$
  - Input-3: MAX resource for each task  $\langle T_{i,j} \rangle$  where  $i=1-N$  and  $j=1-M$ ;
  - Input-4: Sequence of resource request  $\langle R_{i,j} \rangle$  where  $i=1-N$  and  $j=1-M$ 
    - You can define your own way to generate this sequence
  - Test your algorithm with a large number of random sequences of resource request. Make sure deadlock never happens!